

---

# VideoGenDoctor: Auditable Diagnosis and Repair Planning for Controllable Video Generation

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

Controllable video generation pipelines now condition on character identities, camera trajectories, required props, and structured shot scripts. When the output violates these specifications, scalar quality scores from current evaluators cannot localize the failure or recommend a repair. VideoGenDoctor replaces the scalar score with structured diagnosis records. Given a video and a specification, it outputs failure-as-code tuples that each carry a taxonomy code, temporal evidence span, affected target, confidence, representative keyframes, and a compiled repair action.

The default pipeline combines lightweight feature-based screening (CLIP embedding drift, optical flow, face embedding drift, object detection), optional structured VLM verification, and a template-driven patch compiler. On VideoGenDoctor-Bench-v0—a controlled diagnostic fixture of 324 videos spanning 9 perturbation types and 2,016 human-verified evidence spans across 12 observed taxonomy codes—the default diagnosis configuration reaches macro-F1 0.911 and evidence-localization tIoU@0.5 of 0.732. In closed-loop repair under blind human evaluation, VideoGenDoctor-full lifts Pass@2 from 27.8% (score-only) to 76.4%. On a 240-video failure-enriched subset from four production generators, the pipeline retains macro-F1 0.826 and tIoU@0.5 0.598. These results establish that structured diagnosis records with localized evidence enable actionable repair under controlled perturbations. Open-world generalization beyond the current generator and domain coverage remains unverified.

## 1 Introduction

Diffusion-based video generators now accept text, reference images, camera trajectories, character identities, and structured shot scripts as conditioning signals [Yang et al. \[2024\]](#), [Blattmann et al. \[2023\]](#), [Wang et al. \[2023\]](#), [Yuan et al. \[2024\]](#). In a production workflow, a ShotIR-style specification defines per-shot props, characters, camera movement, shot type, and action order. When the output violates these specifications, the user needs fine-grained diagnostics: which failures occurred, where in the timeline, which targets are affected, and what should change in the next generation.

Existing evaluators—VBench [Huang et al. \[2024\]](#), EvalCrafter [Liu et al. \[2024\]](#), VideoScore [He et al. \[2024\]](#), T2V-CompBench [Sun et al. \[2025\]](#)—return scalar quality or alignment scores. These scores rank models but do not localize failures or recommend repairs. A low score offers no signal about whether character identity drifted, a required prop disappeared, camera motion was violated, or a specific segment should be regenerated.

VideoGenDoctor replaces the scalar score with structured diagnosis records. Given a video  $V$  and a specification  $S$ , it outputs

$$R = \{(c_i, t_0^i, t_1^i, y_i, \sigma_i, K_i, T_i, a_i)\},$$

where  $c_i$  is a taxonomy failure code,  $(t_0^i, t_1^i)$  is the temporal evidence span,  $y_i$  names the affected target,  $\sigma_i$  is a confidence score,  $K_i$  provides representative keyframes,  $T_i$  stores optional track-level evidence, and  $a_i$  is a compiled repair action. Each record is both machine-readable and human-auditable.

Three empirical questions follow from this framing. (1) Do structured diagnosis records improve failure-code detection and temporal localization over direct VLM critiques? (2) Do localized evidence and template-constrained repair plans improve downstream video quality beyond score-only feedback? (3) To what extent do conclusions from deterministic perturbations transfer to naturally occurring generator failures?

**Scope.** VideoGenDoctor-Bench-v0 is a controlled diagnostic fixture, not a general open-world benchmark. Its 324 deterministic-perturbation videos isolate known failure modes and support auditable comparison of diagnosis interfaces. The 240-video failure-enriched real-generator subset provides an initial transfer check but cannot estimate natural failure prevalence. Coverage across arbitrary generators, domains, styles, and long-form videos is not established and remains future work.

**Contributions.** This paper introduces: a code–span–target–plan representation for controllable video debugging; a modular diagnosis pipeline that combines feature-based candidate generation, optional structured VLM verification, evidence aggregation, and template-driven patch compilation; VideoGenDoctor-Bench-v0 (324 controlled videos, 2,016 verified evidence spans, 9 perturbation types, 12 observed taxonomy codes, plus a 240-video real-generator subset with 18 observed codes); and an evaluation spanning direct VLM comparisons, blind human repair validation, per-code reliability analysis, paired bootstrap tests, and transfer checks on naturally occurring failures.

## 2 Related work

**Video generation evaluation.** VBench Huang et al. [2024], EvalCrafter Liu et al. [2024], VideoScore He et al. [2024], and T2V-CompBench Sun et al. [2025] measure aggregate quality, motion smoothness, subject consistency, and text-video alignment. They rank models effectively but do not localize failures or produce repair plans. VideoGenDoctor replaces the scalar score with temporally grounded, code-labeled records linked to repair templates (Appendix A, Table 10).

**VLM critics and temporal grounding.** Vision-language models can inspect sampled frames and produce natural-language critiques Liu et al. [2023], Chen et al. [2024], OpenAI [2023]. Direct VLM reporting and VLM-to-patch generation are therefore treated as strong baselines, not strawmen. The key question is whether constraining VLM output through structured code–span–target records improves temporal grounding, schema validity, and downstream repair utility. Temporal alignment is measured by intersection-over-union:


$$\text{tIoU} = \frac{\max(0, \min(\hat{t}_1, t_1^*) - \max(\hat{t}_0, t_0^*))}{\max(\hat{t}_1, t_1^*) - \min(\hat{t}_0, t_0^*)}.$$

**Automated repair.** Program repair localizes defects, classifies them, and generates patches. VideoGenDoctor adapts this pattern to video generation: each diagnosis record exposes evidence and maps to a supported repair family, linking evaluation directly to iterative re-generation.

## 3 VideoGenDoctor

### 3.1 Overview

Figure 1 illustrates the pipeline. It has four stages. (1) The video is partitioned into fixed-stride segments; each segment is sampled into keyframes and four lightweight signals are extracted. (2) Thresholded rules map feature patterns to failure-code candidates with segment-level confidence



**Figure 1:** VideoGenDoctor pipeline: segmentation, feature extraction, rule-based candidate diagnosis, optional VLM verification, patch compilation, and closed-loop re-rendering.

78 scores. (3) An optional VLM judge answers structured yes/no questions on top-ranked candidates  
 79 and produces calibrated confidence scores. (4) A patch compiler maps retained diagnosis records to  
 80 repair actions: reference-frame injection, prop injection, camera override, segment regeneration, or  
 81 style correction. A repair plan is *valid* when it names a supported action, target, and evidence span; it  
 82 is *executable* only when the downstream generator or editor exposes the required control interface.

### 83 3.2 Failure records and patch templates

84 The taxonomy groups 38 failure codes into six categories: Identity, Camera, Motion, Alignment,  
 85 Style, and Scene. The controlled fixture evaluates 12 observed codes; the real-generator subset covers  
 86 additional naturally occurring codes. The remaining codes define an extensible namespace and are  
 87 definitional rather than empirically validated (Appendices A and C). Table 8 in Appendix A lists  
 88 representative failure-to-signal-to-patch mappings.

89 Each predicted failure is a tuple

$$r = (c, t_0, t_1, y, \tilde{\sigma}, K, T, a),$$

90 where  $c$  is the taxonomy code,  $(t_0, t_1)$  spans the evidence,  $y$  names the affected target,  $\tilde{\sigma}$  is the  
 91 confidence,  $K$  holds representative keyframes,  $T$  stores optional track data, and  $a$  is the compiled  
 92 repair action. The schema is detector- and judge-agnostic.

### 93 3.3 Candidate diagnosis and verification

94 The video is split into 2-second non-overlapping segments; six keyframes are sampled per segment.  
 95 Four signals are extracted: semantic embedding drift (OpenCLIP Cherti et al. [2023]), optical flow  
 96 (RAFT Teed and Deng [2020] or Farneback Bradski [2000], Farnebäck [2003]), face embedding drift  
 97 (InsightFace Deng et al. [2019]), and object/prop cues (YOLOv8 Ultralytics [2023], optional). For  
 98 each segment  $s$  and failure code  $c$ , a rule  $\rho_c$  maps the feature vector to a score  $\sigma_{s,c} \in [0, 1]$ . Top-  
 99 ranked candidates inherit the segment’s temporal bounds and the keyframes with highest per-frame  
 100 anomaly scores.

101 The optional Stage-2 judge receives the top- $K$  candidates ( $K = 3$ ), their keyframes, temporal bounds,  
 102 candidate codes, and relevant specification fields. It answers structured yes/no questions rather than  
 103 generating free-form text. The final confidence blends Stage-1 and Stage-2 scores:

$$\tilde{\sigma}_{s,c} = (1 - \alpha)\sigma_{s,c}^{(1)} + \alpha\sigma_{s,c}^{(2)},$$

104 with  $\alpha = 0.6$ . All judge calls log the model, prompt, raw response, latency, and output path.

### 105 3.4 Evidence localization and repair planning

106 Each retained record packs an evidence object  $E = (t_0, t_1, K, T)$  for auditability. The patch compiler  
107 emits either ShotIR field-level diffs (when a structured specification is available) or generator-agnostic  
108 repair plans, preserving the triggering code, evidence span, target, and rationale. In closed-loop mode,  
109 a video is considered fixed when all failure confidences drop below  $\tau = 0.5$ , or after  $K_{\max} = 3$   
110 iterations. Because this stopping criterion depends on the system verifier, Section 4.5 reports  
111 independent blind human validation as the primary metric.

## 112 4 Experiments

### 113 4.1 Benchmark and evaluation protocol

114 VideoGenDoctor is evaluated as a diagnosis-and-repair system, not a video generator. The experiments  
115 target the three questions from the introduction: failure-code correctness, evidence localization, and  
116 repair usefulness under blind human judgment.

117 **Controlled fixture.** VideoGenDoctor-Bench-v0 contains 324 videos built from 18 clean source  
118 clips across six domain groups. Each clip is paired with a ShotIR-style specification and perturbed by  
119 nine deterministic operators under two random seeds, yielding 2,016 human-verified evidence spans  
120 across 12 observed taxonomy codes. The fixture isolates known failure modes to validate interface  
121 consistency, evidence auditability, and repair-plan usefulness under controlled conditions. Construc-  
122 tion details and operator-to-template alignment are in Appendix A (Tables 12) and Appendix B  
123 (Table 20).

124 **Real-failure and real-normal subsets.** We collect 240 naturally occurring failure videos from  
125 CogVideoX, Wan, Stable Video Diffusion, and HunyuanVideo Yang et al. [2024], Wan-AI [2025],  
126 Blattmann et al. [2023], Tencent Hunyuan Foundation Model Team [2025]. From 480 raw genera-  
127 tions, videos are retained only when annotators confirm at least one visible specification violation.  
128 Sampling is balanced at 60 failures per generator, with generator-specific input adapters logged in per-  
129 video manifests: CogVideoX and Wan use ShotIR-to-prompt templates, SVD uses reference-frame  
130 conditioning, and HunyuanVideo uses a multi-shot prompt template. An additional 120 real-normal  
131 videos (30 per generator) are retained from the same raw pool to estimate false positives. Because  
132 the subset is failure-enriched by construction, it supports transfer analysis under verified failures but  
133 cannot estimate natural failure prevalence. Source composition, normal-set behavior, and stress tables  
134 are in Appendix B.

135 **Annotation.** A 120-video reliability subset with 768 candidate evidence spans is dual-annotated  
136 before adjudication. Code-level agreement reaches Cohen’s  $\kappa = 0.831$ ; span-level agreement is  
137 lower (tIoU = 0.681), reflecting inherent ambiguity in fine-grained temporal boundaries. Dataset  
138 statistics and annotation details are in Appendix B (Tables 9–26).

### 139 4.2 Compared methods and metrics

140 We compare three groups of methods. **Diagnosis variants:** Prior-only, Rule-only,  
141 Rule+DummyJudge, Rule+Open-VLM, Rule+GPT-4V, and VideoGenDoctor-diagnosis. **Repair vari-**  
142 **ants:** VideoGenDoctor-patch, Patch+Judge, and VideoGenDoctor-full. **External VLM baselines:**  
143 free-form VLM reporting, structured VLM reporting, self-consistency, temporal proposal + patch,  
144 VLM-to-patch, and Score+LLM-to-patch. Direct VLM baselines share the same GPT-4V endpoint as  
145 Rule+GPT-4V to isolate prompt structure, output schema, and repair protocol from model backbone;  
146 Rule+Open-VLM provides an open-model reference. Diagnosis-only front ends are paired with the  
147 same downstream repair loop whenever human repair success is reported. Implementation details are  
148 in Appendix B (Table 27); strengthened VLM-agent baselines are in Appendix A (Table 13).

149 Metrics include macro-F1 and micro-F1 for failure-code detection; tIoU@0.3/0.5 and Top- $K$   
150 keyframe hit rate for evidence localization; automatic and blind human Pass@1/2 for repair; patch  
151 usefulness (5-point Likert); new-artifact rate; and video-level bootstrap 95% confidence intervals.  
152 Span-aware metrics use Hungarian matching Kuhn [1955] with tIoU as the primary score and confi-  
153 dence as a tie-breaker. Invalid or missing spans receive no localization credit. Paired comparisons

**Table 1:** Failure-code detection on the controlled fixture. VideoGenDoctor-diagnosis outperforms direct VLM reporting; Rule+GPT-4V sets the upper reference. Bootstrap 95% CIs in brackets.

| Method                        | Macro-F1      | Micro-F1      | 95% CI         |
|-------------------------------|---------------|---------------|----------------|
| Prior-only                    | 0.4479        | 0.4635        | [0.419, 0.477] |
| Rule+DummyJudge               | 0.8584        | 0.8716        | [0.836, 0.879] |
| Rule-only                     | 0.8926        | 0.9038        | [0.872, 0.911] |
| Rule+Open-VLM                 | 0.9047        | 0.9175        | [0.885, 0.922] |
| VideoGenDoctor-diagnosis      | 0.9110        | 0.9243        | [0.893, 0.928] |
| Rule+GPT-4V                   | <b>0.9276</b> | <b>0.9386</b> | [0.912, 0.941] |
| VLM free-form report          | 0.8402        | 0.8627        | [0.814, 0.866] |
| VLM structured report         | 0.8869        | 0.9018        | [0.864, 0.908] |
| Human annotator (adjudicated) | 0.958         | 0.967         | [0.942, 0.972] |

**Table 2:** Evidence localization against human-verified spans. Structured diagnosis records improve temporal grounding over free-form VLM outputs. Bootstrap 95% CIs in brackets.

| Method                        | tIoU@0.3      | tIoU@0.5      | Top-1         | Top-3         | 95% CI         |
|-------------------------------|---------------|---------------|---------------|---------------|----------------|
| Prior-only                    | 0.5817        | 0.4295        | 0.1188        | 0.3315        | [0.403, 0.457] |
| Rule+DummyJudge               | 0.7282        | 0.6128        | 0.2241        | 0.4259        | [0.586, 0.639] |
| Rule-only                     | 0.7836        | 0.6714        | 0.2670        | 0.5123        | [0.645, 0.698] |
| Rule+Open-VLM                 | 0.8089        | 0.7062        | 0.3565        | 0.5941        | [0.681, 0.731] |
| VideoGenDoctor-diagnosis      | 0.8261        | 0.7316        | 0.4336        | 0.6815        | [0.707, 0.754] |
| Rule+GPT-4V                   | <b>0.8554</b> | <b>0.7688</b> | <b>0.4923</b> | <b>0.7747</b> | [0.746, 0.790] |
| VLM free-form report          | 0.7204        | 0.5986        | 0.2994        | 0.5117        | [0.570, 0.625] |
| VLM structured report         | 0.7812        | 0.6639        | 0.3651        | 0.5858        | [0.637, 0.690] |
| Human annotator (adjudicated) | 0.882         | 0.815         | 0.531         | 0.762         | [0.790, 0.838] |

report bootstrap differences; small estimates with confidence intervals crossing zero are treated as inconclusive.

### 4.3 Failure-code detection


Table 1 reports failure-code detection results. VideoGenDoctor-diagnosis outperforms both free-form and structured direct VLM reporting on macro-F1; Rule+GPT-4V sets the upper reference. The human adjudicated row demonstrates that even expert annotators do not reach perfect F1 on the controlled fixture. Table 2 reports evidence localization results, showing a corresponding improvement in temporal grounding over VLM baselines.

### 4.4 Per-code reliability and evidence localization

The controlled fixture covers twelve observed codes, with per-code F1 spanning from 0.861 (event-missing) to 0.958 (compression-artifact). Frozen-frame detection reaches the highest F1 (0.956) and tIoU@0.5 (0.794), benefiting from an unambiguous near-zero-flow signal. Event-missing detection and camera-shake detection are hardest to localize (tIoU@0.5 of 0.662 and 0.692), reflecting the inherent ambiguity of temporal boundaries for action absence and high-frequency camera instability. Identity and alignment codes consistently exceed F1 0.90 and tIoU@0.5 of 0.72, while camera and motion codes exhibit greater variance across the group. Table 3 reports full per-code precision, recall, F1, and tIoU. The confusion matrix in Figure 2 confirms that off-diagonal errors are concentrated among camera and motion codes, where overlapping visual artifacts make fine-grained discrimination difficult. Identity, prop-missing, compression, and scene-background failures remain well-separated on the diagonal.

### 4.5 Closed-loop repair and human validation

Table 4 reports closed-loop repair results on the controlled fixture. Automatic Pass@1/2 is computed on all 324 videos using the system verifier, while blind human Pass@1/2 is computed on a stratified



**Figure 2:** Per-code confusion matrix on the controlled fixture. Camera and motion codes show the most off-diagonal confusion; identity, prop, compression, and scene failures are well-separated.

**Table 3:** Per-code reliability of the default VideoGenDoctor diagnosis configuration on the controlled fixture.

| Observed code           | Precision | Recall | F1    | fIoU@0.5 |
|-------------------------|-----------|--------|-------|----------|
| ID_FACE_DRIFT           | 0.938     | 0.914  | 0.926 | 0.748    |
| ID_BODY_DRIFT           | 0.922     | 0.897  | 0.909 | 0.725    |
| CA_MOVE_WRONG           | 0.932     | 0.916  | 0.924 | 0.734    |
| CA_SHOT_TYPE_WRONG      | 0.913     | 0.889  | 0.901 | 0.705    |
| CA_SHAKE                | 0.901     | 0.878  | 0.889 | 0.692    |
| MO_JITTER               | 0.890     | 0.865  | 0.877 | 0.676    |
| MO_FROZEN_FRAME         | 0.963     | 0.950  | 0.956 | 0.794    |
| MO_SEGMENT_BREAK        | 0.906     | 0.883  | 0.894 | 0.704    |
| MO_EVENT_MISSING        | 0.875     | 0.848  | 0.861 | 0.662    |
| AL_PROP_MISSING         | 0.951     | 0.930  | 0.940 | 0.760    |
| ST_COMPRESSION_ARTIFACT | 0.950     | 0.966  | 0.958 | 0.779    |
| SC_BG_INCONSISTENCY     | 0.910     | 0.885  | 0.897 | 0.686    |

177 144-video subset. Human raters see the specification and output video but not method names or  
178 diagnosis records, so the human metric is the primary check against verifier-coupled gains.

179 Random-patch and re-render-all baselines demonstrate that undirected repair is substantially  
180 weaker than diagnosis-guided patching: re-render-all achieves Pass@2 of 0.543 (vs. 0.764 for  
181 VideoGenDoctor-full) at higher regeneration cost. The paired comparison in Section 4.7 shows a  
182 small and non-decisive difference between Patch+Judge and VideoGenDoctor-full.



**Table 4:** Closed-loop repair on the controlled fixture. Automatic Pass@K on 324 videos; blind human Pass@K on a stratified 144-video subset. VideoGenDoctor-full lifts human Pass@2 from 27.8% (score-only) to 76.4%.

| Method              | Auto P@1      | 95% CI         | Auto P@2      | 95% CI         | Human Pass@1  | 95% CI         | Human Pass@2  | 95% CI         | Patch useful | New artifacts |
|---------------------|---------------|----------------|---------------|----------------|---------------|----------------|---------------|----------------|--------------|---------------|
| Random-patch        | 0.148         | [0.113, 0.187] | 0.296         | [0.249, 0.344] | 0.118         | [0.075, 0.172] | 0.236         | [0.175, 0.303] | 1.80         | 0.194         |
| Re-render-all       | 0.296         | [0.248, 0.347] | 0.543         | [0.489, 0.596] | 0.264         | [0.197, 0.336] | 0.500         | [0.419, 0.579] | 2.61         | 0.153         |
| Score-only          | 0.2407        | [0.196, 0.291] | 0.3580        | [0.307, 0.411] | 0.1875        | [0.130, 0.257] | 0.2778        | [0.207, 0.355] | 2.28         | 0.0833        |
| Patch-only          | 0.5247        | [0.471, 0.579] | 0.6975        | [0.646, 0.744] | 0.4792        | [0.399, 0.559] | 0.6319        | [0.551, 0.706] | 3.50         | 0.1736        |
| VLM-to-patch        | 0.5617        | [0.507, 0.615] | 0.7438        | [0.694, 0.788] | 0.5069        | [0.427, 0.586] | 0.6736        | [0.594, 0.746] | 3.64         | 0.2153        |
| Patch+Judge         | 0.6296        | [0.575, 0.681] | 0.8179        | [0.773, 0.857] | 0.5625        | [0.482, 0.640] | 0.7431        | [0.667, 0.807] | 3.93         | 0.1458        |
| VideoGenDoctor-full | <b>0.6481</b> | [0.594, 0.699] | <b>0.8364</b> | [0.793, 0.873] | <b>0.5833</b> | [0.503, 0.660] | <b>0.7639</b> | [0.689, 0.826] | <b>4.11</b>  | 0.1319        |

**Table 5:** Controlled-to-real transfer. All methods degrade on natural failures; the gap between controlled and real subsets quantifies the difficulty of uncurated artifacts.

| Subset              | Method                              | Macro-F1      | tIoU@0.5      | Patch useful | Human Pass@2  |
|---------------------|-------------------------------------|---------------|---------------|--------------|---------------|
| Controlled fixture  | VLM structured report + repair loop | 0.8869        | 0.6639        | 3.62         | 0.6810        |
| Controlled fixture  | VideoGenDoctor default pipeline     | <b>0.9110</b> | <b>0.7316</b> | <b>4.11</b>  | <b>0.7639</b> |
| Real-failure subset | VLM structured report + repair loop | 0.7992        | 0.5431        | 3.36         | 0.6075        |
| Real-failure subset | VideoGenDoctor default pipeline     | <b>0.8264</b> | <b>0.5982</b> | <b>3.82</b>  | <b>0.6908</b> |

## 4.6 Controlled fixture versus real generator failures

Table 5 compares performance on the controlled and real-failure subsets. Figure 4 in Appendix A visualizes closed-loop repair; Figure 6 in Appendix B shows the transfer in a paired bar-chart format.

All methods degrade on naturally occurring failures, confirming that the controlled fixture is not a substitute for open-world evaluation. Because this subset is failure-enriched and balanced by construction, these numbers measure transfer under verified failures rather than real-world prevalence. Per-code breakdowns on the real-failure subset, per-generator results, and the paired bar-chart comparison are in Appendix B (Tables 28–29, Figure 6).

## 4.7 Ablation and generalization checks

**Feature signal contribution.** We isolate each feature signal by removing it from the default configuration and by using each signal alone. Table 6 reports the results. Optical flow is the single most critical signal: removing it drops macro-F1 from 0.9110 to 0.849 and tIoU@0.5 from 0.7316 to 0.651, consistent with its role in detecting motion and camera failures. CLIP embedding drift is the second most important signal, particularly for scene and style codes. Face embedding removal mainly degrades identity-code precision, while object detection has the smallest individual impact because YOLOv8 is active only for prop-related codes. Single-signal configurations (CLIP only, flow only, face only) all underperform substantially, confirming that the signals carry complementary information. Figure 7 in Appendix B provides a horizontal bar-chart comparison.

**Repair component ablation.** Table 7 isolates the contribution of each diagnosis component to downstream repair. Each record component adds measurable repair value beyond code prediction alone. The jump from code-only (Pass@2 0.458) to code+span (0.625) and code+span+evidence (0.688) shows that temporal evidence is the largest single contributor after the code itself. Full visualizations are in Appendix B (Figure 8).

We further examine sensitivity to three hyperparameters. First, sweeping the diagnosis retention threshold over  $\tau \in \{0.30, \dots, 0.70\}$  shows that macro-F1 and tIoU@0.5 peak at  $\tau = 0.50$ ; lower thresholds over-trigger repairs, while higher thresholds suppress useful diagnoses (Appendix B, Table 15, Figure 9). Second, the Stage-2 judge weight  $\alpha = 0.6$  and candidate budget  $K = 3$  jointly give the best operating point. Removing Stage-2 ( $\alpha = 0$ ) reduces both diagnosis and repair metrics, while  $\alpha = 0.9$  introduces instability from VLM over-correction (Appendix B, Figure 10). Third, paired bootstrap comparisons confirm that verification adds a reliable gain (+0.111 in Human Pass@2), while the full-loop wrapper contributes a small, non-decisive increment (+0.021) (Appendix B, Table 16).

Finally, leave-one-perturbation-out and leave-one-source-group-out evaluations confirm that performance drops under held-out operators—particularly missing-event, camera-shake, and temporal-discontinuity perturbations—but remains above prior-only and score-only baselines (Appendix B, Tables 30–31). The system does not simply memorize perturbation templates, though unseen-operator

**Table 6:** Per-feature ablation of the default diagnosis configuration. Each row removes one signal or uses a single signal; Human Pass@2 is measured after pairing with the shared repair loop.

| Configuration      | Macro-F1 | tIoU@0.5 | Human Pass@2 | $\Delta$ Pass@2 |
|--------------------|----------|----------|--------------|-----------------|
| Full (all signals) | 0.9110   | 0.7316   | 0.7639       | —               |
| — CLIP embedding   | 0.871    | 0.683    | 0.704        | −0.060          |
| — Optical flow     | 0.849    | 0.651    | 0.672        | −0.092          |
| — Face embedding   | 0.876    | 0.695    | 0.718        | −0.046          |
| — Object detection | 0.893    | 0.708    | 0.737        | −0.027          |
| CLIP only          | 0.762    | 0.581    | 0.583        | −0.181          |
| Flow only          | 0.638    | 0.472    | 0.465        | −0.299          |
| Face only          | 0.524    | 0.385    | 0.398        | −0.366          |

**Table 7:** Repair ablation isolating the additive contribution of failure codes, target fields, evidence spans, keyframes, patch templates, and verification. Human Pass@2 on the 144-video blind subset.

| Repair setting           | Code     | Target   | Span     | Evidence | Template | Human Pass@2  | Patch useful | New artifacts |
|--------------------------|----------|----------|----------|----------|----------|---------------|--------------|---------------|
| Score-only               | ✗        | ✗        | ✗        | ✗        | ✗        | 0.2778        | 2.28         | 0.0833        |
| Code-only patch          | ✓        | ✗        | ✗        | ✗        | ✓        | 0.4583        | 2.91         | 0.1597        |
| Code+target patch        | ✓        | ✓        | ✗        | ✗        | ✓        | 0.5417        | 3.24         | 0.1736        |
| Code+span patch          | ✓        | ✗        | ✓        | ✗        | ✓        | 0.6250        | 3.57         | 0.1667        |
| Code+span+evidence patch | ✓        | ✗        | ✓        | ✓        | ✓        | 0.6875        | 3.80         | 0.1528        |
| VLM-to-patch             | optional | optional | optional | optional | ✗        | 0.6736        | 3.64         | 0.2153        |
| Patch+Judge              | ✓        | ✓        | ✓        | ✓        | ✓        | 0.7431        | 3.93         | 0.1458        |
| VideoGenDoctor-full      | ✓        | ✓        | ✓        | ✓        | ✓        | <b>0.7639</b> | <b>4.11</b>  | 0.1319        |

generalization is weaker than in-distribution evaluation. Additional checks on adapter executability, real-failure stress slices, and multi-annotator stability are in Appendix B (Tables 17–19).

## 5 Discussion

The evidence supports a bounded claim: structured diagnosis records improve failure detection, temporal grounding, and downstream repair over score-only feedback and direct VLM reporting on the current evaluation sets. The controlled fixture isolates known perturbations for auditable comparison; the real-failure subset provides an initial transfer check that confirms the approach degrades under natural artifacts but remains useful.

VideoGenDoctor is best understood as an interface contribution, not a new generator or learned repair algorithm. Its advantages are schema validity, evidence auditability, repair-plan structure, and cost-awareness. Rule+GPT-4V achieves the strongest absolute scores but at higher cost and latency; VideoGenDoctor-full is the recommended default because it preserves the full audit trail at the paper’s reference cost point. The modular design allows feature extractors, rules, VLM judges, and patch compilers to be replaced independently.

**Limitations.** The controlled fixture does not represent open-world video diversity. The real-failure subset is failure-enriched—it supports transfer analysis under verified failures but cannot estimate natural prevalence. Automatic pass rates are partly verifier-dependent; we report blind human evaluation as the primary repair metric. Only the 12 observed taxonomy codes carry empirical support; the remaining 38 codes are definitional. Repair assumes a downstream generator or editor that consumes at least part of the patch vocabulary. The annotation study uses two adjudicated annotators; broader validation requires more generators, longer videos, additional domains, and larger annotator pools.

**Verifier–human agreement.** Because the closed-loop stopping criterion depends on the system verifier, we measure verifier–human agreement on the 144-video blind-evaluation subset. Agreement between the automatic verifier and the majority human judgment reaches 0.841 (Cohen’s  $\kappa = 0.786$ ), indicating that the automatic pass/fail signal is well-calibrated but not a perfect substitute for human evaluation. The primary Pass@K claims in this paper are therefore based on blind human judgment; automatic Pass@K serves as a scalable proxy for ablation and hyperparameter sweeps.

**Break-even analysis.** At the default operating point, diagnosis costs approximately \$0.24 per video (CLIP + flow + face features, rule engine, and Stage-2 VLM on  $K = 3$  candidates). Full



re-rendering costs approximately \$0.52 per video on current hardware. The break-even failure rate is approximately 46%: when fewer than half of generated videos contain specification violations, diagnosis-guided repair is cheaper than re-rendering everything. When failure rates are higher—as in early-stage prompt engineering or new-domain adaptation—re-render-all may be more economical. This trade-off is visualized in Appendix A, Figure 5.

**When does it fail?** Diagnosis degrades most on overlapping artifacts where multiple failure codes are plausible (e.g., camera shake vs. motion jitter), on subtle style shifts below the CLIP embedding sensitivity threshold, and on failures in the last 20% of the video timeline where fewer keyframes are sampled. The Stage-2 VLM judge occasionally over-corrects when asked about codes outside its effective visual range (e.g., fine-grained shot-type discrimination). False positives on real-normal videos are concentrated among texture-flicker and lighting-shift codes where the rule engine triggers on benign high-frequency variation. These failure modes inform the recommendations for threshold selection and candidate budget in Section 4.7.

## 6 Conclusion

VideoGenDoctor replaces scalar video evaluation with auditable diagnosis records that each locate a failure in time, name the affected target, and prescribe a repair action. On a 324-video controlled diagnostic fixture, the default pipeline detects failure codes at macro-F1 0.911, localizes evidence spans at tIoU@0.5 of 0.732, and lifts blind human repair Pass@2 from 27.8% to 76.4%. On a 240-video failure-enriched real-generator subset spanning four production generators and 18 taxonomy codes, the pipeline retains macro-F1 0.826 and tIoU@0.5 0.598. The modular design—separating feature extraction, rule-based screening, VLM verification, and patch compilation—allows each component to be replaced independently. These results support structured diagnosis-and-repair as a practical interface for controllable video generation, but do not establish generality beyond the controlled perturbations, generator set, and video lengths evaluated here.

## 7 Ethics statement

VideoGenDoctor improves the controllability of generated video through automated diagnosis. Improved generation reliability could be misused to strengthen misleading media. In sensitive deployments, automated repair should be paired with provenance tracking, watermarking, and human review. The failure taxonomy encodes normative judgments about visual quality; creative workflows should expose configurable thresholds and permit human override.

## References

- Andreas Blattmann et al. Stable video diffusion: Scaling latent video diffusion models to large datasets, 2023. URL <https://arxiv.org/abs/2311.15127>. arXiv preprint arXiv:2311.15127.
- Gary Bradski. The opencv library. *Dr. Dobb’s Journal of Software Tools*, 2000.
- Zhe Chen et al. InternVL: Scaling up vision foundation models and aligning for generic visual-linguistic tasks. *CVPR*, 2024.
- Mehdi Cherti et al. Reproducible scaling laws for contrastive language-image learning. *CVPR*, 2023.
- Jiankang Deng, Jia Guo, Niannan Xue, and Stefanos Zafeiriou. ArcFace: Additive angular margin loss for deep face recognition. *CVPR*, 2019.
- Gunnar Farnebäck. Two-frame motion estimation based on polynomial expansion. In *Scandinavian Conference on Image Analysis*, pages 363–370. Springer, 2003.
- Xuan He et al. VideoScore: Building automatic metrics to simulate fine-grained human feedback for video generation. *ACL*, 2024.
- Ziqi Huang, Yinan He, Jiashuo Yu, et al. VBench: Comprehensive benchmark suite for video generative models. *CVPR*, 2024.
- Harold W. Kuhn. The hungarian method for the assignment problem. *Naval Research Logistics Quarterly*, 2(1–2):83–97, 1955. doi: 10.1002/nav.3800020109.

**Table 8:** Failure types, the evidence signals that detect them, and the repair actions they trigger.

| Failure type                   | Evidence signal                                    | Repair action                                 |
|--------------------------------|----------------------------------------------------|-----------------------------------------------|
| Face identity drift            | Face embedding drift across keyframes              | Reference-frame injection                     |
| Body / clothing drift          | Character-region appearance drift                  | Reference-frame or character-anchor injection |
| Camera movement deviation      | Camera flow trace or trajectory mismatch           | Camera movement override                      |
| Frozen frames or action stalls | Near-zero flow or stalled pose track               | Regenerate affected segment                   |
| Missing required prop          | Object presence check or VLM prop-absence judgment | Prop injection                                |
| Style inconsistency            | Style or color embedding drift across segments     | Style correction                              |
| Background inconsistency       | Scene/background embedding drift                   | Background anchoring                          |

**Table 9:** Dataset statistics for the controlled fixture and failure-enriched real-generator subset. The real-failure split cannot estimate natural failure prevalence.

| Split               | #Vid | #Evidence | AvgDur | PerturbTypes | #Codes |
|---------------------|------|-----------|--------|--------------|--------|
| Controlled fixture  | 324  | 2016      | 10.46s | 9            | 12     |
| Real-failure subset | 240  | 1536      | 8.93s  | –            | 18     |

- 296 Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. *NeurIPS*, 2023.
- 297 Yaofang Liu et al. EvalCrafter: Benchmarking and evaluating large video generation models. *CVPR*, 2024.
- 298 OpenAI. GPT-4V(ision) system card. Technical report, OpenAI, 2023. URL <https://openai.com/research/gpt-4v-system-card>.
- 299 Kaiyue Sun, Kaiyi Huang, Xian Liu, Yue Wu, Zihan Xu, Zhenguo Li, and Xihui Liu. T2V-CompBench: A comprehensive benchmark for compositional text-to-video generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 8406–8416, 2025.
- 300 Zachary Teed and Jia Deng. RAFT: Recurrent all-pairs field transforms for optical flow. *ECCV*, 2020.
- 304 Tencent Hunyuan Foundation Model Team. Hunyuanvideo 1.5 technical report, 2025. URL <https://arxiv.org/abs/2511.18870>.
- 305 Ultralytics. YOLOv8. Software repository, 2023. URL <https://github.com/ultralytics/ultralytics>.
- 306 Version 8.0.
- 308 Wan-AI. Wan2.1-t2v-14b. Hugging Face model card, 2025. URL <https://huggingface.co/Wan-AI/Wan2.1-T2V-14B>.
- 310 Zhouxia Wang, Ziyang Yuan, Xintao Wang, Tianshui Chen, Menghan Xia, Ping Luo, and Ying Shan. MotionCtrl: A unified and flexible motion controller for video generation, 2023. URL <https://arxiv.org/abs/2312.03641>. arXiv preprint arXiv:2312.03641.
- 312 Zhuoyi Yang et al. CogVideoX: Text-to-video diffusion models with an expert transformer, 2024. URL <https://arxiv.org/abs/2408.06072>. arXiv preprint arXiv:2408.06072.
- 313 Hangjie Yuan, Shiwei Zhang, Xiang Wang, Yujie Wei, Tao Feng, Yining Pan, Yingya Zhang, Ziwei Liu, Samuel Albanie, and Dong Ni. InstructVideo: Instructing video diffusion models with human feedback. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 6463–6474, 2024.

## 319 A Supplementary main-text results

320 This appendix stores tables and figures moved out of the main text to keep the core narrative compact  
 321 while preserving the audit trail for comparisons, taxonomy scope, fixture construction, temporal  
 322 evidence profiles, strengthened VLM baselines, repair visualization, and cost-performance trade-offs.

**Table 10:** Capability comparison with related evaluation and critique paradigms. Here “Global score” denotes scalar quality, alignment, or verifier-style pass outputs rather than localized evidence records. Direct VLM reporting is treated as a strong baseline rather than only comparing against score-only benchmarks.

| System                | Global score | Code    | Evidence | Patch   | Closed-Loop |
|-----------------------|--------------|---------|----------|---------|-------------|
| VBench                | ✓            | ✗       | ✗        | ✗       | ✗           |
| EvalCrafter           | ✓            | ✗       | ✗        | ✗       | ✗           |
| VideoScore            | ✓            | ✗       | ✗        | ✗       | ✗           |
| T2V-CompBench         | ✓            | partial | ✗        | ✗       | ✗           |
| Direct VLM report     | partial      | partial | partial  | partial | ✗           |
| <b>VideoGenDoctor</b> | ✓            | ✓       | ✓        | ✓       | ✓           |

**Table 11:** Failure-as-code taxonomy groups. Each group defines a namespace for diagnosis records and links to evidence procedures and repair-plan templates.

| Group     | Operational scope                                                                 |
|-----------|-----------------------------------------------------------------------------------|
| Identity  | Identity drift, face/body consistency, wrong character, clothing consistency.     |
| Camera    | Shot type, camera movement, viewpoint angle, unintended zoom, camera shake.       |
| Motion    | Action timing, frame continuity, motion physics, frozen frames, segment breaks.   |
| Alignment | Prompt adherence, required objects and props, character count, prop placement.    |
| Style     | Color consistency, texture flicker, compression artifacts, art-style consistency. |
| Scene     | Background consistency, lighting, environment match, scene-object stability.      |

## B Additional experimental tables

This appendix collects construction details, audit tables, stress analyses, per-slice results, and visualizations that support the main experimental claims without interrupting the main narrative.

### B.1 Diagnosis and repair visualizations

### B.2 Extended experiment analyses

**Adapter executability.** Table 17 and Figure 11 separate repair-plan quality from backend executability by grouping generated plans into L0–L3 adapter levels. The trend is consistent with the intended repair interface: L2/L3 actions are more often executable and yield higher Human Pass@2, while L0/L1 outputs remain useful as conservative abstentions or prompt-level repair instructions but provide weaker direct repair success.

**Stronger real-failure stress slices.** Table 18 and Figure 12 extend the real-failure check to harder slices: additional generators, longer videos, style-shifted prompts, and more complex prompt specifications. The stress slices degrade relative to the base real-failure subset, especially for long videos and complex prompts, but the retained performance remains above the score-only and prior-only operating ranges reported in the main experiments.

**Multi-annotator stability.** Table 19 and Figure 13 extend the annotation reliability analysis from dual annotation to a three-annotator setting. Code-level agreement remains substantial, while span agreement is lower on real-failure samples, consistent with the higher ambiguity of naturally occurring failures and less isolated temporal evidence.

## C Full failure taxonomy

Table 32 reports the machine-readable taxonomy used by VideoGenDoctor v0.1. The v0.1 taxonomy contains 38 codes grouped into six categories. Each code is paired with an operational definition, an evidence procedure, and a default patch template. The current benchmark evaluates the 12 codes observed in VideoGenDoctor-Bench-v0, while the remaining codes define the extensible namespace used by the report schema and patch compiler.

**Table 12:** Composition of the controlled diagnostic fixture. Each domain group contains three clean source clips. Every source clip is paired with a ShotIR-style specification and perturbed by nine deterministic operators under two random seeds.

| Domain group               | #Source clips | Avg. duration | Resolution | #Perturb. ops | #Seeds | #Videos |
|----------------------------|---------------|---------------|------------|---------------|--------|---------|
| Indoor / human-object      | 3             | 10.1s         | 768×432    | 9             | 2      | 54      |
| Outdoor / motion           | 3             | 10.8s         | 768×432    | 9             | 2      | 54      |
| Multi-object interaction   | 3             | 10.4s         | 768×432    | 9             | 2      | 54      |
| Camera-control scene       | 3             | 11.0s         | 768×432    | 9             | 2      | 54      |
| Style / lighting variation | 3             | 10.2s         | 768×432    | 9             | 2      | 54      |
| Scene-change / background  | 3             | 10.3s         | 768×432    | 9             | 2      | 54      |
| Total                      | 18            | 10.46s        | –          | 9             | 2      | 324     |

**Table 13:** Strengthened VLM-agent baselines on the controlled fixture. All rows are evaluated under the same downstream repair protocol; diagnosis-only front ends are explicitly paired with that shared repair loop when human repair success is reported. Self-consistency improves direct VLM reporting but increases cost. The validity column reports whether the emitted action stays inside the supported repair-plan vocabulary; it does not claim generator-adaptor executability. Best values are boldfaced; higher is better except for relative cost, where lower is better.


| Method                            | Macro-F1      | tIoU@0.5      | Human Pass@2  | Schema validity | Patch vocabulary validity | Relative cost |
|-----------------------------------|---------------|---------------|---------------|-----------------|---------------------------|---------------|
| VLM free-form report              | 0.8402        | 0.5986        | 0.6125        | 0.742           | 0.604                     | 1.12×         |
| VLM structured report             | 0.8869        | 0.6639        | 0.6810        | 0.914           | 0.732                     | 1.16×         |
| VLM structured + self-consistency | 0.8978        | 0.6824        | 0.7042        | 0.951           | 0.755                     | 3.48×         |
| VLM temporal proposal + patch     | 0.8915        | 0.7018        | 0.7076        | 0.936           | 0.802                     | 2.05×         |
| VideoGenDoctor-full               | 0.9110        | 0.7316        | 0.7639        | 0.991           | 0.953                     | <b>1.00×</b>  |
| Rule+GPT-4V + repair loop         | <b>0.9276</b> | <b>0.7688</b> | <b>0.7917</b> | <b>0.993</b>    | <b>0.957</b>              | 2.75×         |

**Table 32:** Full VideoGenDoctor v0.1 failure taxonomy used by the report schema and patch compiler.

| Code                    | Definition                                                         | Evidence Procedure                                                        | Patch Template                    |
|-------------------------|--------------------------------------------------------------------|---------------------------------------------------------------------------|-----------------------------------|
| Identity                |                                                                    |                                                                           |                                   |
| ID_FACE_DRIFT           | Face identity changes inconsistently between frames or segments.   | Compare face embeddings across keyframes and flag large cosine distance.  | inject_<br>reference_<br>frame    |
| ID_BODY_DRIFT           | Body appearance, clothing, or build changes inconsistently.        | Compare character-region embeddings across segments.                      | inject_<br>reference_<br>frame    |
| ID_WRONG_CHARACTER      | A visible character does not match any specified character.        | Check face embeddings against known character anchors.                    | replace_<br>character_<br>anchor  |
| ID_CLOTHING_CHANGE      | Clothing changes between shots without narrative justification.    | Measure character-region appearance drift across segments.                | fix_<br>character_<br>consistency |
| ID_MULTIPLE_FACES       | Multiple distinct faces appear when one character is expected.     | Count distinct identities detected in a single-character shot.            | adjust_<br>character_<br>count    |
| ID_NO_FACE_VISIBLE      | An expected face is absent when the shot requires face visibility. | Check face detection under close-up or medium-shot constraints.           | adjust_<br>framing                |
| Scene                   |                                                                    |                                                                           |                                   |
| SC_BG_INCONSISTENCY     | The background changes unexpectedly between segments.              | Measure background-region embedding drift across consecutive segments.    | fix_<br>background_<br>anchor     |
| SC_ENVIRONMENT_MISMATCH | The scene environment does not match the specified location.       | Compare frames with the location description using image-text similarity. | set_<br>environment               |
| SC_LIGHTING_SHIFT       | Lighting changes inconsistently within a shot.                     | Measure per-frame luminance variation within the segment.                 | normalize_<br>lighting            |
| SC_OBJECT_TELEPORT      | A scene object appears or disappears without plausible motion.     | Track object boxes across adjacent frames.                                | stabilize_<br>scene_objects       |
| SC_BG_FLICKER           | The background flickers or pulses unnaturally.                     | Detect high-frequency luminance variation in background regions.          | fix_<br>background_<br>anchor     |
| SC_WRONG_TIME_OF_DAY    | The time of day does not match the specification.                  | Classify sky and ambient lighting against the expected time.              | set_time_of_<br>day               |
| Motion                  |                                                                    |                                                                           |                                   |

Continued on next page

| Code                       | Definition                                                              | Evidence Procedure                                                       | Patch Template                     |
|----------------------------|-------------------------------------------------------------------------|--------------------------------------------------------------------------|------------------------------------|
| MO_JITTER                  | Unnatural frame-to-frame jitter appears outside intended camera motion. | Measure optical-flow magnitude variance and direction instability.       | normalize_<br>frame_rate           |
| MO_FRAME_DROP              | Duplicate or dropped frames produce visible stutter.                    | Detect near-zero adjacent-frame flow when motion is expected.            | interpolate_<br>frames             |
| MO_UNNATURAL_MOTION        | Character or object motion is physically implausible.                   | Compare flow direction and magnitude with expected motion constraints.   | constrain_<br>motion               |
| MO_EVENT_MISSING           | A required action or event does not occur in the expected segment.      | Ask the VLM judge whether the specified action is visible.               | inject_<br>action_<br>keyframe     |
| MO_ACTION_ORDER_WRONG      | Actions occur in the wrong temporal order.                              | Compare detected action timestamps with the specification order.         | reorder_<br>segments               |
| MO_MOTION_BLUR_EXCESS      | Excessive motion blur degrades keyframes.                               | Use sharpness measures such as Laplacian variance.                       | reduce_<br>motion_blur             |
| MO_FROZEN_FRAME            | The video is frozen when motion is expected.                            | Detect near-zero flow across the full segment.                           | regenerate_<br>segment             |
| MO_SEGMENT_BREAK           | A visual discontinuity occurs at a segment boundary.                    | Detect embedding-drift spikes at boundaries.                             | smooth_<br>segment_<br>transition  |
| Camera                     |                                                                         |                                                                          |                                    |
| CA_MOVE_WRONG              | Camera movement deviates from the specified movement.                   | Compare optical-flow direction patterns with the expected movement type. | set_camera_<br>movement            |
| CA_SHOT_TYPE_WRONG         | The shot type does not match the specification.                         | Compare face or body box size ratios with the expected shot type.        | adjust_<br>framing                 |
| CA_UNINTENDED_ZOOM         | Unplanned zoom-in or zoom-out occurs.                                   | Detect radial flow without a zoom specification.                         | remove_zoom                        |
| CA_SHAKE                   | Camera shake appears outside the specified movement.                    | Detect high-frequency oscillation in global optical flow.                | stabilize_<br>camera               |
| CA_WRONG_ANGLE             | Camera angle deviates from the specification.                           | Ask the VLM judge to compare the observed angle with the target angle.   | set_camera_<br>angle               |
| CA_ROLL                    | Unintended camera roll appears.                                         | Estimate the rotational component of optical flow.                       | remove_<br>camera_roll             |
| Alignment                  |                                                                         |                                                                          |                                    |
| AL_PROP_MISSING            | A required prop is not visible in the segment.                          | Use object detection or VLM verification for prop absence.               | inject_prop                        |
| AL_PROP_WRONG_PLACEMENT    | A required prop appears in the wrong region.                            | Compare the prop box with the expected screen region.                    | reposition_<br>prop                |
| AL_TEXT_PROMPT_MISMATCH    | The generated content does not align with the prompt.                   | Measure text-image similarity and verify mismatches with a judge.        | strengthen_<br>prompt_<br>guidance |
| AL_CHARACTER_COUNT_WRONG   | The number of visible characters does not match the specification.      | Count faces or bodies and compare with the expected count.               | adjust_<br>character_<br>count     |
| AL_WRONG_PROP              | A visible prop is not listed in the specification.                      | Detect high-confidence unlisted objects.                                 | remove_prop                        |
| AL_PROP_OCCLUDED           | A required prop is present but substantially occluded.                  | Estimate visible area and occlusion of the prop box.                     | reposition_<br>prop                |
| Style                      |                                                                         |                                                                          |                                    |
| ST_COLOR_SHIFT             | The color palette changes inconsistently across segments.               | Measure HSV histogram distance between segments.                         | apply_color_<br>grading            |
| ST_ART_STYLE_INCONSISTENCY | The art style shifts between segments.                                  | Compare style embeddings across segments.                                | apply_style_<br>transfer           |
| ST_COMPRESSION_ARTIFACT    | Visible compression artifacts degrade video quality.                    | Detect blocking, ringing, or quality-score degradation.                  | apply_<br>enhancement              |
| ST_TEXTURE_FLICKER         | Texture patterns flicker or shimmer unnaturally.                        | Measure high temporal frequency in texture regions.                      | smooth_<br>texture                 |
| ST_OVEREXPOSURE            | Significant frame regions are overexposed.                              | Count the fraction of saturated high-value pixels.                       | adjust_<br>exposure                |
| ST_UNDEREXPOSURE           | Significant frame regions are underexposed.                             | Count the fraction of near-black pixels.                                 | adjust_<br>exposure                |



**Figure 3:** Temporal distribution of verified evidence spans in the evaluation sets. Panel (a) aggregates the 2,016 controlled-fixture spans over 2-second bins and shows the front- and mid-loaded pattern expected from deterministic perturbation operators. Panel (b) summarizes the 1,536 real-failure spans over the same time bins, making the later and more diffuse temporal mass directly comparable to panel (a). Panel (c) further breaks the real-failure spans down by failure group and time bin, highlighting that motion and camera failures dominate the later intervals and overlap more strongly than in the controlled fixture.


## D Judge protocol

The Stage-2 judge receives the top- $K$  candidate segments produced by Stage-1. For each segment, the input includes the segment identifier, temporal bounds, keyframe paths, candidate failure codes, and specification context, including required props, expected actions, camera movement, shot type, and character anchors. The judge answers structured questions rather than producing a free-form report. This design keeps the output comparable across local open-source VLMs and hosted vision-capable models.

For identity failures, the judge checks whether the face, body appearance, clothing, or other distinctive visual attributes of the main character change inconsistently across keyframes. For alignment failures, it verifies whether required props are visible, whether they appear in the expected screen region, and whether the generated content matches the prompt or structured specification. For camera failures, it compares the observed camera movement, shot type, and viewpoint angle with the requested control



Closed-loop repair comparison



**Figure 4:** Closed-loop repair comparison on the controlled fixture. Left: automatic verifier Pass@1/2 on all 324 controlled videos. Right: blind human Pass@1/2 on the stratified 144-video human-evaluation subset. Automatic pass rates may be coupled to the system verifier, while human pass rates provide an independent check of repair success.

**Table 14:** Cost-performance trade-off under a shared downstream repair protocol. Diagnosis-only front ends are paired with the same repair loop before reporting Human Pass@2. Cost is reported relative to the default VideoGenDoctor-full configuration. Values in the last two columns are multiplicative factors relative to VideoGenDoctor-full (1.00×). Best values are boldfaced; higher is better for task metrics, lower is better for cost and latency.

| Method                      | Macro-F1      | tIoU@0.5      | Human Pass@2  | Relative cost | Relative latency |
|-----------------------------|---------------|---------------|---------------|---------------|------------------|
| Rule-only + repair loop     | 0.8926        | 0.6714        | 0.5988        | <b>0.42×</b>  | <b>0.58×</b>     |
| Rule+Open-VLM + repair loop | 0.9047        | 0.7062        | 0.7160        | 0.73×         | 0.87×            |
| VideoGenDoctor-full         | 0.9110        | 0.7316        | 0.7639        | 1.00×         | 1.00×            |
| Rule+GPT-4V + repair loop   | <b>0.9276</b> | <b>0.7688</b> | <b>0.7917</b> | 2.75×         | 2.31×            |
| VLM structured report       | 0.8869        | 0.6639        | 0.6810        | 1.16×         | 1.20×            |

360 signal. For motion failures, it verifies whether required actions occur, whether they occur in the  
 361 specified order, and whether stutter, frozen frames, or segment breaks are visible. For style and scene  
 362 failures, it checks lighting, color, compression artifacts, background consistency, and environment  
 363 match.

364 The output is a structured object containing a natural-language answer for each question, a confidence  
 365 score, per-code probability updates, and a re-ranking of keyframes. The final failure score is computed  
 366 as

$$\tilde{\sigma}_{s,c} = (1 - \alpha)\sigma_{s,c}^{(1)} + \alpha\sigma_{s,c}^{(2)},$$


367 where  $\sigma_{s,c}^{(1)}$  is the Stage-1 score,  $\sigma_{s,c}^{(2)}$  is the judge score, and  $\alpha = 0.6$  by default. All judge calls  
 368 record the model name, prompt, raw response, token count when available, and output path to support  
 369 reproducibility.

## 370 E Patch template examples

371 VideoGenDoctor compiles each diagnosis record into one or more patch actions. If a ShotIR  
 372 specification is available, the compiler emits field-level diffs. If the generator does not expose a  
 373 structured specification interface, the compiler emits generator-agnostic repair plans that can be  
 374 translated into prompt edits, reference-frame injection, segment-level re-rendering, or post-processing  
 375 steps.

```

376 {"action": "inject_reference_frame",
377   "field": "character_id",
378   "value": "<anchor_frame>",
379   "reason": "ID_FACE_DRIFT detected at t=1.2s"}
```



**Figure 5:** Pareto view of the cost-performance trade-off under the shared downstream repair protocol. The default VideoGenDoctor-full operating point improves human repair success over lower-cost rule variants and direct structured VLM reporting, while the higher-cost Rule+GPT-4V reference obtains the strongest Human Pass@2 at substantially higher relative cost.

**Table 15:** Threshold sensitivity.  $\tau = 0.50$  balances diagnosis coverage against false positives and unnecessary patching.

| Threshold $\tau$ | Macro-F1     | tIoU@0.5     | Real-normal FPR | Unnecessary patch rate |
|------------------|--------------|--------------|-----------------|------------------------|
| 0.30             | 0.902        | 0.716        | 0.158           | 0.142                  |
| 0.40             | 0.910        | 0.728        | 0.125           | 0.111                  |
| 0.50             | <b>0.911</b> | <b>0.732</b> | 0.100           | 0.087                  |
| 0.60             | 0.904        | 0.719        | 0.083           | 0.071                  |
| 0.70             | 0.886        | 0.694        | <b>0.067</b>    | <b>0.058</b>           |

```

380 {"action": "set_camera_movement",
381      "field": "camera_movement",
382      "value": "<expected_camera_movement>",
383      "reason": "CA_MOVE_WRONG detected in the localized span"}

```

```

384 {"action": "inject_prop",
385      "field": "props_required",
386      "value": "<missing_prop>",
387      "reason": "AL_PROP_MISSING verified by object or VLM evidence"}

```


```

388 {"action": "regenerate_segment",
389      "field": "segment",
390      "value": "<t0,t1>",
391      "reason": "MO_FROZEN_FRAME or MO_SEGMENT_BREAK is localized to this span"}


```

Each patch specifies an action type, the affected target, an optional update value, and a short rationale tied to the evidence span. When multiple records affect the same target, the compiler groups them before re-rendering to reduce conflicts among patch actions.

In the reported Human Pass@K evaluation, the concretely executed actions are the L2/L3 subset exposed by the active backend, primarily reference-frame injection, segment rerendering or temporal



**Figure 6:** Controlled-to-real transfer across four metrics. Solid bars: controlled fixture; hatched bars: real-failure subset. The degradation gap is largest for tIoU@0.5 and smallest for patch usefulness.




**Figure 7:** Per-feature ablation horizontal bar chart. Optical flow is the single most critical signal; single-signal configurations all underperform substantially.


397 smoothing, enhancement or deblocking, and mask- or reference-conditioned prop insertion when  
 398 supported. L0/L1 rows remain structured repair-plan instructions and are not counted as executable  
 399 patches.

## 400 F Annotation guide summary

401 Annotators review each perturbed video together with the original prompt or ShotIR specification  
 402 and the auto-assigned candidate failure codes. For each candidate, annotators determine whether  
 403 the failure is visually present, assign the tightest temporal span that contains the visible deviation,  
 404 and select representative keyframes when available. Annotators may also add missing failures if the  
 405 perturbation produces an additional visible artifact.




**Figure 8:** Repair ablation as grouped bars. Human Pass@2 and patch usefulness rise additively as code, span, evidence, and verification are incorporated.



**Figure 9:** Threshold sensitivity curves. Dashed line at  $\tau = 0.50$  marks the operating point where macro-F1 and tIoU@0.5 peak.

The annotation record is stored as one JSON object per video and includes the video identifier, annotator identifier, verified failure codes, confidence values, evidence spans, and free-form notes. A failure is retained in the benchmark only when the annotator can verify it from the rendered video and specification, without relying on perturbation metadata alone. Ambiguous cases are handled conservatively: if a deviation is not visually identifiable, the candidate is marked as not verified.

The experiment fixture includes a dual-annotated reliability subset of 120 videos and 768 candidate evidence spans. The subset contains both controlled perturbation samples and naturally occurring real-failure samples. Larger multi-annotator validation with more than two independent annotators remains future work.



**Figure 10:** Sensitivity to judge weight  $\alpha$  and candidate budget  $K$ .  $\alpha = 0.6$ ,  $K = 3$  balances diagnosis accuracy and repair success.

**Table 16:** Paired bootstrap differences in blind human Pass@2. Verification adds a reliable gain; the full-loop wrapper adds a small, non-decisive increment.

| Comparison                         | $\Delta$ Human Pass@2 | 95% CI          | Interpretation       |
|------------------------------------|-----------------------|-----------------|----------------------|
| Patch-only vs Score-only           | 0.3541                | [0.250, 0.458]  | substantial gain     |
| VLM-to-patch vs Patch-only         | 0.0417                | [-0.054, 0.139] | small / not decisive |
| Patch+Judge vs Patch-only          | 0.1112                | [0.021, 0.215]  | verification helps   |
| VideoGenDoctor-full vs Patch+Judge | 0.0208                | [-0.063, 0.104] | small / not decisive |

## 415 G Example diagnostic report

## 416 H Direct VLM baseline prompt templates


417 The four direct VLM/LLM baselines differ only in their allowed information access and output schema.  
 418 The free-form and structured VLM report baselines receive sampled frames and the specification.  
 419 The VLM-to-patch baseline also receives sampled frames, the taxonomy names, and the repair-plan  
 420 vocabulary, but not VideoGenDoctor’s rule scores or patch compiler outputs. The Score+LLM-to-  
 421 patch baseline receives scalar evaluator scores and the specification only; it does not receive sampled  
 422 frames. These information boundaries are fixed before evaluation and are used for all samples.

423 **F.1 Free-form VLM report prompt.** The model receives the video specification, sampled  
 424 keyframes with timestamps, and a short instruction: identify visible failures, explain the evidence,  
 425 and propose repair suggestions. The response is then parsed into taxonomy codes and evidence spans  
 426 for evaluation.

427 **F.2 Structured VLM report prompt.** The model receives the same input but must output  
 428 JSON objects with fields `failure_code`, `temporal_span`, `affected_target`, `confidence`,  
 429 `evidence_keyframes`, `rationale`, and `repair_action`. Failures without visual evidence must  
 430 be marked as uncertain rather than hallucinated.

**Table 17:** Adapter-executability stratification. L0: abstention; L1: prompt-level instructions; L2: partially executable actions; L3: native generator/editor controls.

| Adapter level | Plan share | Adapter-executable rate | Human Pass@2 | New artifacts |
|---------------|------------|-------------------------|--------------|---------------|
| L0 abstain    | 0.100      | 0.000                   | 0.280        | 0.040         |
| L1 plan       | 0.310      | 0.280                   | 0.540        | 0.100         |
| L2 partial    | 0.420      | 0.740                   | 0.730        | 0.140         |
| L3 native     | 0.170      | 0.960                   | 0.810        | 0.160         |



**Figure 11:** Adapter-executability stratification by repair-plan level. The left panel reports the share of generated plans by adapter level; the right panel compares adapter-executable rate, Human Pass@2, and new-artifact rate within each level.


431 You are evaluating whether a generated video follows the given specification.  
432 For each visible failure, return a JSON record:  
433 {  
434   "failure\_code": "<taxonomy code>",  
435   "temporal\_span": [t0, t1],  
436   "affected\_target": "<character / prop / camera / motion / style / scene>",  
437   "confidence": <0-1>,  
438   "evidence\_keyframes": ["frame@time", ...],  
439   "rationale": "short visual reason",  
440   "repair\_action": "<patch action>"  
441 }  
442 Only report failures that are visually supported by the frames.

443 **F.3 VLM-to-patch prompt.** The VLM-to-patch baseline receives the same video specification and  
444 sampled keyframes as the direct VLM diagnosis baselines, but it is asked to output repair actions  
445 directly rather than diagnosis records. This baseline is allowed to access the taxonomy names and the  
446 supported repair-plan vocabulary, but it is not given VideoGenDoctor’s Stage-1 rule scores, candidate  
447 segments, or patch compiler outputs. It may infer temporal spans from the sampled frames, but it  
448 does not receive ground-truth spans or VideoGenDoctor-predicted spans. This setting tests whether a  
449 VLM can directly produce structured repair plans from the specification and visual evidence.

450 You are a video-generation repair planner.  
451  
452 Input:  
453 1. A video specification or ShotIR-style instruction.  
454 2. Sampled keyframes from the generated video, each with a timestamp.  
455 3. The failure-code taxonomy names.  
456 4. The supported repair-plan vocabulary.  
457  
458 Your task:  
459 Inspect the sampled frames and the specification. Identify only visually supported

**Table 18:** Stress validation on harder real-failure slices. Intervals are bootstrap half-widths for Figure 12.

| Slice             | Macro-F1          | tIoU@0.5          | Human Pass@2      | FPR               |
|-------------------|-------------------|-------------------|-------------------|-------------------|
| Base real-failure | 0.826 $\pm$ 0.020 | 0.598 $\pm$ 0.017 | 0.691 $\pm$ 0.028 | 0.100 $\pm$ 0.010 |
| New generators    | 0.802 $\pm$ 0.024 | 0.574 $\pm$ 0.020 | 0.662 $\pm$ 0.034 | 0.126 $\pm$ 0.012 |
| Long videos       | 0.784 $\pm$ 0.029 | 0.552 $\pm$ 0.025 | 0.628 $\pm$ 0.041 | 0.142 $\pm$ 0.014 |
| Style shift       | 0.811 $\pm$ 0.025 | 0.581 $\pm$ 0.022 | 0.674 $\pm$ 0.036 | 0.118 $\pm$ 0.013 |
| Complex prompts   | 0.793 $\pm$ 0.027 | 0.560 $\pm$ 0.023 | 0.641 $\pm$ 0.039 | 0.135 $\pm$ 0.014 |




**Figure 12:** Stress validation on harder real-failure slices. Left: diagnosis, localization, and human repair with uncertainty. Right: false-positive rate on matched real-normal controls.

460 problems and output structured repair-plan actions. Do not write a free-form critique.  
461 Do not invent failures that are not visible in the frames.  
462  
463 You may use the following failure-code groups:  
464 - Identity: ID\_FACE\_DRIFT, ID\_BODY\_DRIFT, ID\_WRONG\_CHARACTER,  
465 ID\_CLOTHING\_CHANGE, ID\_MULTIPLE\_FACES, ID\_NO\_FACE\_VISIBLE  
466 - Camera: CA\_MOVE\_WRONG, CA\_SHOT\_TYPE\_WRONG, CA\_UNINTENDED\_ZOOM,  
467 CA\_SHAKE, CA\_WRONG\_ANGLE, CA\_ROLL  
468 - Motion: MO\_JITTER, MO\_FRAME\_DROP, MO\_UNNATURAL\_MOTION,  
469 MO\_EVENT\_MISSING, MO\_ACTION\_ORDER\_WRONG, MO\_MOTION\_BLUR\_EXCESS,  
470 MO\_FROZEN\_FRAME, MO\_SEGMENT\_BREAK  
471 - Alignment: AL\_PROP\_MISSING, AL\_PROP\_WRONG\_PLACEMENT,  
472 AL\_TEXT\_PROMPT\_MISMATCH, AL\_CHARACTER\_COUNT\_WRONG,  
473 AL\_WRONG\_PROP, AL\_PROP\_OCCLUDED  
474 - Style: ST\_COLOR\_SHIFT, ST\_ART\_STYLE\_INCONSISTENCY,  
475 ST\_COMPRESSION\_ARTIFACT, ST\_TEXTURE\_FLICKER,  
476 ST\_OVEREXPOSURE, ST\_UNDEREXPOSURE  
477 - Scene: SC\_BG\_INCONSISTENCY, SC\_ENVIRONMENT\_MISMATCH,  
478 SC\_LIGHTING\_SHIFT, SC\_OBJECT\_TELEPORT, SC\_BG\_FLICKER,  
479 SC\_WRONG\_TIME\_OF\_DAY  
480  
481 Supported repair-plan vocabulary:  
482 - inject\_reference\_frame  
483 - replace\_character\_anchor  
484 - fix\_character\_consistency  
485 - adjust\_framing  
486 - set\_camera\_movement  
487 - set\_camera\_angle  
488 - stabilize\_camera  
489 - remove\_zoom  
490 - remove\_camera\_roll  
491 - inject\_prop  
492 - reposition\_prop



**Table 19:** Three-annotator stability analysis. Boundary disagreement in seconds (lower is better).

| Subset       | #Videos | #Candidate spans | Annotators | Fleiss' $\kappa$ | Mean pairwise tIoU / Boundary disagreement |
|--------------|---------|------------------|------------|------------------|--------------------------------------------|
| Controlled   | 72      | 468              | 3          | 0.842            | 0.704 / 0.38s                              |
| Real-failure | 48      | 300              | 3          | 0.782            | 0.628 / 0.52s                              |
| Combined     | 120     | 768              | 3          | 0.813            | 0.666 / 0.45s                              |

**Figure 13:** Multi-annotator stability. Code-level  $\kappa$  exceeds temporal-boundary agreement; the real-failure split shows highest ambiguity.

```

493 - remove_prop
494 - adjust_character_count
495 - strengthen_prompt_guidance
496 - inject_action_keyframe
497 - reorder_segments
498 - regenerate_segment
499 - interpolate_frames
500 - normalize_frame_rate
501 - constrain_motion
502 - smooth_segment_transition
503 - reduce_motion_blur
504 - apply_color_grading
505 - apply_style_transfer
506 - apply_enhancement
507 - smooth_texture
508 - adjust_exposure
509 - fix_background_anchor
510 - set_environment
511 - normalize_lighting
512 - stabilize_scene_objects
513 - set_time_of_day
514 - no_op_uncertain
515
516 Return a JSON object with the following schema:
517 {
518   "patches": [
519     {
520       "patch_id": "<unique patch id>",
521       "failure_code": "<taxonomy code or uncertain>",
522       "action": "<one item from the supported repair-plan vocabulary>",
523       "target": "<character / prop / camera / motion / style / scene target>",
524       "temporal_span": [<start_sec>, <end_sec>],
525       "evidence_keyframes": ["frame@time", "..."],
526       "confidence": <0-1>,

```

**Table 20:** Operator-template alignment analysis for the controlled fixture. The table makes explicit which observable cues, rule signals, failure codes, and repair-plan templates are aligned by construction.

| Perturbation operator                  | Observable cue                                             | Rule / feature signal                                  | Failure code                 | Repair-plan template                             | Alignment type      |
|----------------------------------------|------------------------------------------------------------|--------------------------------------------------------|------------------------------|--------------------------------------------------|---------------------|
| Identity-anchor removal                | face or body appearance drift                              | face embedding drift; character-region embedding drift | ID_FACE_DRIFT, ID_BODY_DRIFT | inject_reference_frame; replace_character_anchor | identity-anchor     |
| Required-prop dropping                 | required object becomes absent or inconsistent             | object detector absence; VLM prop-absence judgment     | AL_PROP_MISSING              | inject_prop                                      | prop-presence       |
| Camera-movement alteration             | requested trajectory is replaced by static or wrong motion | global flow direction; trajectory mismatch             | CA_MOVE_WRONG                | set_camera_movement                              | motion-control      |
| Shot-type / framing alteration         | framing violates requested shot type                       | crop statistics; VLM framing judgment                  | CA_SHOT_TYPE_WRONG           | adjust_framing                                   | framing-control     |
| Camera-shake injection                 | unstable camera with unintended jitter                     | flow instability; camera-shake heuristic               | CA_SHAKE                     | stabilize_camera                                 | stability-control   |
| Action dropping / shifting             | required event missing or misordered                       | action-track gap; VLM event-order judgment             | MO_EVENT_MISSING             | inject_action_keyframe; regenerate_segment       | event-structure     |
| Temporal jitter / frame-drop injection | discontinuity, duplicate frames, or rate mismatch          | flow anomaly; near-zero motion plateau                 | MO_JITTER, MO_SEGMENT_BREAK  | normalize_frame_rate; regenerate_segment         | temporal-continuity |
| Compression re-encoding                | blockiness and de-blocking artifacts                       | compression detector; texture degradation signal       | ST_COMPRESSION_ARTIFACT      | apply_enhancement                                | quality-restoration |
| Background-anchor weakening            | scene background drifts across the shot                    | scene embedding drift; background-anchor mismatch      | SC_BG_INCONSISTENCY          | fix_background_anchor                            | scene-anchor        |

**Table 21:** Source composition of the failure-enriched real-generator subset. We generate 480 raw videos and retain 240 samples with annotator-verified visible failures. Sampling is balanced at 60 retained failures per generator, and generator-specific prompt or conditioning adapters are logged in the construction manifest. The subset is used only as an initial transfer check, not to estimate natural failure prevalence. We use publicly identifiable generator checkpoints to make the subset construction auditable and reproducible.

| Generator              | Version / checkpoint | Input type                       | #Prompts / specs | #Videos | Avg. duration | #Unique observed codes |
|------------------------|----------------------|----------------------------------|------------------|---------|---------------|------------------------|
| CogVideoX              | CogVideoX1.5-5B      | text / ShotIR-to-prompt          | 60               | 60      | 8.4s          | 15                     |
| Wan                    | Wan2.1-T2V-14B       | text / ShotIR-to-prompt          | 60               | 60      | 9.2s          | 16                     |
| Stable Video Diffusion | SVD-XT-1.1           | image-to-video / reference frame | 60               | 60      | 8.6s          | 14                     |
| HunyuanVideo           | HunyuanVideo-1.5     | text / multi-shot prompt         | 60               | 60      | 9.5s          | 17                     |
| Total                  | –                    | –                                | 240              | 240     | 8.93s         | 18                     |

```

527     "patch_fields": {
528         "field": "<generator or ShotIR field to modify>",
529         "value": "<new value, constraint, or reference to insert>"
530     },
531     "rationale": "<short visual reason>"
532 }
533 ],
534 "global_notes": "<one-sentence summary>",
535 "uncertain": <true or false>
536 }
537
538 Rules:
539 1. Output valid JSON only.
540 2. Every patch must use an action from the supported repair-plan vocabulary.
541 3. If no visually supported failure is visible, return:
542     {"patches": [], "global_notes": "No visually supported repair action.", "uncertain": true}
543 4. If the temporal span is uncertain, estimate it from the nearest sampled frame
544     timestamps and set "confidence" below 0.5.
545 5. Do not output patches for purely aesthetic preferences unless the specification
546     explicitly requires them.

```

**F.4 Score+LLM-to-patch prompt.** The Score+LLM-to-patch baseline receives scalar evaluator outputs and the original video specification, but it does not receive sampled frames. This baseline is intentionally weaker in visual grounding but represents a common production setting where low-level evaluators return scalar scores and an LLM converts low-scoring dimensions into repair suggestions.

**Table 22:** Real-failure and real-normal evaluation. The real-normal subset contains 120 videos retained from the same 480 raw generations but judged to contain no obvious failure. Whenever a diagnosis-oriented front end is listed, patch-related columns are computed after pairing that front end with the shared downstream repair loop. Best values are boldfaced; higher is better except for real-normal FPR and unnecessary patch rate, where lower is better.

| Method                            | Real-failure Recall | Real-normal FPR | Unnecessary Patch Rate | No-op Accuracy | Specificity  |
|-----------------------------------|---------------------|-----------------|------------------------|----------------|--------------|
| Score-only                        | 0.621               | 0.258           | 0.231                  | 0.769          | 0.742        |
| Rule-only                         | 0.806               | 0.168           | 0.154                  | 0.846          | 0.832        |
| VLM structured report             | 0.818               | 0.217           | 0.205                  | 0.795          | 0.783        |
| VLM structured + self-consistency | 0.837               | 0.167           | 0.150                  | 0.850          | 0.833        |
| VLM temporal proposal + patch     | 0.848               | 0.192           | 0.183                  | 0.817          | 0.808        |
| VideoGenDoctor-diagnosis          | <b>0.862</b>        | 0.108           | 0.092                  | 0.908          | 0.892        |
| VideoGenDoctor-full               | 0.858               | <b>0.100</b>    | <b>0.087</b>           | <b>0.913</b>   | <b>0.900</b> |

**Table 23:** Stress analysis on 72 ambiguous specification-violation cases. These samples contain possible but non-adjudicable deviations from the specification. Whenever a diagnosis-oriented front end is listed, patch-related columns are computed after pairing that front end with the shared downstream repair loop. Best values are boldfaced; higher is better for abstention-oriented columns, lower is better for false concrete patch, wrong-target patch, over-repair, and new artifacts.

| Method                        | Abstention rate | Appropriate abstention | False concrete patch | Wrong-target patch | Over-repair  | New artifacts |
|-------------------------------|-----------------|------------------------|----------------------|--------------------|--------------|---------------|
| Score-only                    | 0.292           | 0.361                  | 0.514                | 0.222              | 0.347        | 0.083         |
| VLM structured report         | 0.181           | 0.292                  | 0.611                | 0.278              | 0.431        | 0.153         |
| VLM temporal proposal + patch | 0.222           | 0.347                  | 0.556                | 0.236              | 0.389        | 0.194         |
| VideoGenDoctor-diagnosis      | <b>0.500</b>    | <b>0.639</b>           | 0.292                | 0.125              | 0.208        | <b>0.069</b>  |
| VideoGenDoctor-full           | 0.472           | 0.625                  | <b>0.250</b>         | <b>0.111</b>       | <b>0.181</b> | 0.111         |

551 Because it cannot inspect frames, it is not allowed to claim frame-specific visual evidence. It may  
552 output approximate temporal spans only when the scalar evaluator provides segment-level scores;  
553 otherwise the temporal span must be marked as null. This makes the information boundary explicit  
554 and prevents unfair comparison with frame-based VLM baselines.

555 You are a repair planner for controllable video generation.

556

557 Input:

- 558 1. The original video specification or ShotIR-style instruction.
- 559 2. Scalar evaluator outputs, such as quality, alignment, motion, identity,  
560 camera, style, and scene-consistency scores.
- 561 3. Optional segment-level scalar scores if available.
- 562 4. The supported repair-plan vocabulary.

563

564 Important information boundary:

565 You do NOT receive sampled frames or the video itself. Therefore:

- 566 - Do not claim that a failure is visually observed.
- 567 - Do not mention specific visual evidence or keyframes.
- 568 - Only infer likely repair actions from low-scoring evaluator dimensions.
- 569 - Use "evidence\_source": "scalar\_score" for all patches.
- 570 - Use "temporal\_span": null unless segment-level scores identify a specific interval.

571

572 Supported repair-plan vocabulary:

- 573 - inject\_reference\_frame
- 574 - replace\_character\_anchor
- 575 - fix\_character\_consistency
- 576 - adjust\_framing
- 577 - set\_camera\_movement
- 578 - set\_camera\_angle
- 579 - stabilize\_camera
- 580 - remove\_zoom
- 581 - remove\_camera\_roll
- 582 - inject\_prop
- 583 - reposition\_prop
- 584 - remove\_prop

**Table 24:** Stress analysis on 48 low-quality or non-adjudicable cases. These samples are dominated by global degradation, unclear content, or insufficient evidence for reliable code-level labels. Whenever a diagnosis-oriented front end is listed, patch-related columns are computed after pairing that front end with the shared downstream repair loop. Best values are boldfaced; higher is better for quality-gated abstention, lower is better for the remaining columns.

| Method                        | Quality-gated abstention | False concrete patch | Wrong-target patch | Over-repair  | New artifacts |
|-------------------------------|--------------------------|----------------------|--------------------|--------------|---------------|
| Score-only                    | 0.417                    | 0.375                | 0.125              | 0.208        | 0.063         |
| VLM structured report         | 0.292                    | 0.458                | 0.188              | 0.271        | 0.104         |
| VLM temporal proposal + patch | 0.333                    | 0.396                | 0.167              | 0.250        | 0.146         |
| VideoGenDoctor-diagnosis      | <b>0.667</b>             | 0.146                | 0.063              | 0.104        | <b>0.042</b>  |
| VideoGenDoctor-full           | 0.625                    | <b>0.125</b>         | <b>0.042</b>       | <b>0.083</b> | 0.083         |

**Table 25:** Observed failure-code distribution in the controlled fixture and real-failure subset. The full taxonomy contains 38 codes, but empirical claims are restricted to the observed codes below.

| Failure code            | Group     | Controlled spans | Real-failure spans | Total spans |
|-------------------------|-----------|------------------|--------------------|-------------|
| ID_FACE_DRIFT           | Identity  | 150              | 112                | 262         |
| ID_BODY_DRIFT           | Identity  | 132              | 76                 | 208         |
| ID_CLOTHING_CHANGE      | Identity  | –                | 60                 | 60          |
| CA_MOVE_WRONG           | Camera    | 240              | 104                | 344         |
| CA_SHOT_TYPE_WRONG      | Camera    | 144              | 84                 | 228         |
| CA_SHAKE                | Camera    | 126              | 70                 | 196         |
| CA_WRONG_ANGLE          | Camera    | –                | 68                 | 68          |
| MO_JITTER               | Motion    | 174              | 124                | 298         |
| MO_FROZEN_FRAME         | Motion    | 156              | 72                 | 228         |
| MO_SEGMENT_BREAK        | Motion    | 162              | 88                 | 250         |
| MO_EVENT_MISSING        | Motion    | 126              | 90                 | 216         |
| MO_ACTION_ORDER_WRONG   | Motion    | –                | 56                 | 56          |
| AL_PROP_MISSING         | Alignment | 198              | 116                | 314         |
| AL_TEXT_PROMPT_MISMATCH | Alignment | –                | 132                | 132         |
| ST_COMPRESSION_ARTIFACT | Style     | 252              | 80                 | 332         |
| ST_COLOR_SHIFT          | Style     | –                | 56                 | 56          |
| SC_BG_INCONSISTENCY     | Scene     | 156              | 86                 | 242         |
| SC_OBJECT_TELEPORT      | Scene     | –                | 62                 | 62          |
| Total                   | –         | 2016             | 1536               | 3552        |

```

585 - adjust_character_count
586 - strengthen_prompt_guidance
587 - inject_action_keyframe
588 - reorder_segments
589 - regenerate_segment
590 - interpolate_frames
591 - normalize_frame_rate
592 - constrain_motion
593 - smooth_segment_transition
594 - reduce_motion_blur
595 - apply_color_grading
596 - apply_style_transfer
597 - apply_enhancement
598 - smooth_texture
599 - adjust_exposure
600 - fix_background_anchor
601 - set_environment
602 - normalize_lighting
603 - stabilize_scene_objects
604 - set_time_of_day
605 - no_op_uncertain

```

```

606
607 Return a JSON object with the following schema:
608 {
609   "patches": [
610     {
611       "patch_id": "<unique patch id>",

```

**Table 26:** Annotation reliability. Dual annotation is performed on a reliability subset before adjudication.

| Subset                          | #Videos | #Candidate spans | #Annotators | Code agreement   | Span agreement |
|---------------------------------|---------|------------------|-------------|------------------|----------------|
| Controlled reliability subset   | 72      | 468              | 2           | $\kappa = 0.858$ | tIoU = 0.716   |
| Real-failure reliability subset | 48      | 300              | 2           | $\kappa = 0.803$ | tIoU = 0.641   |
| Combined                        | 120     | 768              | 2           | $\kappa = 0.831$ | tIoU = 0.681   |

**Table 27:** Implementation details for VLM-based judges and external diagnosis/repair baselines.

| Setting                           | Model                  | Frames             | Res.             | Prompt               | Rel. cost |
|-----------------------------------|------------------------|--------------------|------------------|----------------------|-----------|
| Rule+Open-VLM                     | Qwen2-VL-7B-Instruct   | 18                 | 512px short side | structured QA        | 0.34×     |
| Rule+GPT-4V                       | GPT-4V endpoint        | 18                 | 512px short side | structured QA        | 2.75×     |
| VLM free-form report              | GPT-4V endpoint        | 12                 | 512px short side | free-form critique   | 1.12×     |
| VLM structured report             | GPT-4V endpoint        | 12                 | 512px short side | JSON schema          | 1.16×     |
| VLM structured + self-consistency | GPT-4V endpoint        | 12 frames × 3 runs | 512px short side | JSON schema, 3 runs  | 3.48×     |
| VLM temporal proposal + patch     | GPT-4V endpoint        | 12                 | 512px short side | span proposal + JSON | 2.05×     |
| VLM-to-patch                      | GPT-4V endpoint        | 12                 | 512px short side | JSON patch           | 1.19×     |
| Score+LLM-to-patch                | scalar evaluator + LLM | 0                  | –                | score-to-patch JSON  | N/A       |

*Note.* Most direct VLM baselines share the same hosted GPT-4V endpoint [OpenAI \[2023\]](#) to isolate prompt- and schema-level effects rather than backbone changes; Rule+Open-VLM provides an open-model reference. Hosted endpoint calls are logged with raw responses for auditability. ‘N/A’ denotes a scalar-evaluator-dominated pipeline whose cost is not directly comparable to frame-based VLM calls.

```

612     "score_dimension": "<low-scoring evaluator dimension>",
613     "inferred_failure_type": "<short inferred failure description>",
614     "action": "<one item from the supported repair-plan vocabulary>",
615     "target": "<character / prop / camera / motion / style / scene target>",
616     "temporal_span": null,
617     "evidence_source": "scalar_score",
618     "triggering_scores": {
619         "<score_name>": <score_value>
620     },
621     "confidence": <0-1>,
622     "patch_fields": {
623         "field": "<generator or ShotIR field to modify>",
624         "value": "<new value, constraint, or reference to insert>"
625     },
626     "rationale": "<short reason based only on scalar scores and specification>"
627 }
628 ],
629 "global_notes": "<one-sentence summary>",
630 "uncertain": <true or false>
631 }
632
633 Rules:
634 1. Output valid JSON only.
635 2. Do not claim visual evidence because frames are not provided.
636 3. Do not output frame IDs or evidence keyframes.
637 4. Use temporal_span = null unless segment-level scalar scores are provided.
638 5. Every patch must use an action from the supported repair-plan vocabulary.
639 6. If the scalar scores do not indicate a clear repair target, return:
640     {"patches": [], "global_notes": "No reliable score-based repair action.", "uncertain": true}

```

**F.5 Parsing and invalid-output handling.** Invalid JSON outputs are repaired using a deterministic parser when possible. If the output cannot be parsed, the prediction is marked invalid. Failure codes outside the taxonomy are mapped to the nearest taxonomy code only when an exact synonym exists; otherwise they are counted as false positives. Predictions without temporal spans receive no localization credit. Hallucinated failures without visual evidence are counted as false positives. Repair actions outside the supported repair-plan vocabulary are marked invalid for repair-plan scoring and non-adaptor-executable for adaptor scoring.

**Table 28:** Per-code reliability on the real-failure subset. The subset is more challenging than the controlled fixture, especially for motion, camera, and scene-level failures.

| Failure code            | #Spans | Precision | Recall | F1    | tIoU@0.5 |
|-------------------------|--------|-----------|--------|-------|----------|
| ID_FACE_DRIFT           | 112    | 0.884     | 0.848  | 0.866 | 0.648    |
| ID_BODY_DRIFT           | 76     | 0.856     | 0.823  | 0.839 | 0.615    |
| ID_CLOTHING_CHANGE      | 60     | 0.834     | 0.798  | 0.816 | 0.586    |
| CA_MOVE_WRONG           | 104    | 0.866     | 0.832  | 0.849 | 0.628    |
| CA_SHOT_TYPE_WRONG      | 84     | 0.838     | 0.803  | 0.820 | 0.590    |
| CA_SHAKE                | 70     | 0.820     | 0.779  | 0.799 | 0.559    |
| CA_WRONG_ANGLE          | 68     | 0.825     | 0.786  | 0.805 | 0.568    |
| MO_JITTER               | 124    | 0.823     | 0.777  | 0.799 | 0.565    |
| MO_FROZEN_FRAME         | 72     | 0.899     | 0.854  | 0.876 | 0.663    |
| MO_SEGMENT_BREAK        | 88     | 0.827     | 0.792  | 0.809 | 0.582    |
| MO_EVENT_MISSING        | 90     | 0.807     | 0.762  | 0.784 | 0.542    |
| MO_ACTION_ORDER_WRONG   | 56     | 0.792     | 0.748  | 0.769 | 0.524    |
| AL_PROP_MISSING         | 116    | 0.893     | 0.863  | 0.878 | 0.669    |
| AL_TEXT_PROMPT_MISMATCH | 132    | 0.842     | 0.794  | 0.817 | 0.596    |
| ST_COMPRESSION_ARTIFACT | 80     | 0.914     | 0.874  | 0.893 | 0.688    |
| ST_COLOR_SHIFT          | 56     | 0.846     | 0.806  | 0.825 | 0.601    |
| SC_BG_INCONSISTENCY     | 86     | 0.832     | 0.795  | 0.813 | 0.578    |
| SC_OBJECT_TELEPORT      | 62     | 0.823     | 0.782  | 0.802 | 0.545    |

**Table 29:** Per-generator results on the real-failure and real-normal subsets. FPR is computed on 30 real-normal samples per generator.

| Generator              | Macro-F1 | tIoU@0.5 | Human Pass@2 | Real-normal FPR | Main failure groups       | Normal samples |
|------------------------|----------|----------|--------------|-----------------|---------------------------|----------------|
| CogVideoX              | 0.834    | 0.609    | 0.702        | 0.092           | Camera, Motion, Alignment | 30             |
| Wan                    | 0.821    | 0.588    | 0.681        | 0.108           | Motion, Camera, Identity  | 30             |
| Stable Video Diffusion | 0.842    | 0.617    | 0.696        | 0.083           | Identity, Style, Scene    | 30             |
| HunyuanVideo           | 0.809    | 0.579    | 0.684        | 0.117           | Camera, Motion, Alignment | 30             |
| Average                | 0.827    | 0.598    | 0.691        | 0.100           | —                         | 120            |

## I Blind human repair evaluation protocol

Raters are shown the specification and one output video, without method names or diagnosis records. They answer whether the video satisfies the specification overall, whether the target failure is resolved, whether the repair introduces new artifacts, and how useful the patch is on a 1–5 scale. Human Pass@ $K$  is the fraction of samples judged successful after at most  $K$  repair attempts.

A repaired video is marked as pass if the target failure is resolved and no severe new artifact is introduced. Human Pass@ $K$  is the fraction of samples judged successful after at most  $K$  repair attempts. New artifacts are marked when raters observe a visible degradation introduced by repair that was not present in the original generated video. Each repaired video is rated by 3 independent raters. We report the majority-vote pass/fail label, the mean patch-usefulness score, and the majority-vote new-artifact label.

## J Bootstrap confidence intervals

We use 10,000 video-level bootstrap resamples and report percentile 95% confidence intervals. For the controlled fixture, resampling is performed over 324 videos. For the real-failure subset, resampling is performed over 240 videos. For blind human repair evaluation, resampling is performed over the stratified human-rated subset. For paired comparisons, we resample videos with replacement and compute the distribution of metric differences between two methods. Small point-estimate differences are not treated as meaningful unless the paired confidence interval excludes zero.

The large degradation under leave-one-code-family-out ablation confirms that VideoGenDoctor should not be interpreted as an unconstrained open-ended failure discovery system. The experiment clarifies which portions of the performance depend on family-specific detectors and which portions remain supported by generic evidence signals and VLM verification.

**Table 30:** Leave-one-perturbation-out evaluation. Each held-out operator is excluded from threshold tuning and then used only for testing.

| Held-out perturbation operator         | Macro-F1 | tIoU@0.5 | Top-3 hit |
|----------------------------------------|----------|----------|-----------|
| Identity-anchor removal                | 0.884    | 0.681    | 0.644     |
| Required-prop dropping                 | 0.902    | 0.704    | 0.672     |
| Camera-movement alteration             | 0.871    | 0.653    | 0.612     |
| Shot-type / framing alteration         | 0.862    | 0.641    | 0.606     |
| Camera-shake injection                 | 0.851    | 0.628    | 0.590     |
| Action dropping / shifting             | 0.838    | 0.612    | 0.568     |
| Temporal jitter / frame-drop injection | 0.846    | 0.621    | 0.577     |
| Compression re-encoding                | 0.910    | 0.719    | 0.688     |
| Background-anchor weakening            | 0.868    | 0.646    | 0.610     |
| Average                                | 0.870    | 0.656    | 0.619     |

**Table 31:** Leave-one-source-group-out evaluation over the six controlled-fixture domain groups.

| Held-out source group      | Macro-F1 | tIoU@0.5 | Top-3 hit |
|----------------------------|----------|----------|-----------|
| Indoor / human-object      | 0.899    | 0.714    | 0.665     |
| Outdoor / motion           | 0.887    | 0.702    | 0.651     |
| Multi-object interaction   | 0.892    | 0.706    | 0.657     |
| Camera-control scene       | 0.875    | 0.683    | 0.631     |
| Style / lighting variation | 0.907    | 0.722    | 0.676     |
| Scene-change / background  | 0.884    | 0.694    | 0.642     |
| Average                    | 0.891    | 0.704    | 0.654     |

## K Real-failure construction manifest

When a generator does not expose one of these configuration fields, the corresponding manifest value is set to null rather than omitted. For hosted or version-changing endpoints, we log the endpoint identifier, access date, raw request payload, and raw response metadata.

To make the real-failure subset auditable and reproducible, each generated video is accompanied by a JSON manifest. The manifest records the generator identity, checkpoint, input specification, generation configuration, output location, screening decision, verified failure codes, temporal evidence spans, and annotator identifiers. This protocol turns the real-failure subset from an informal collection of generated examples into a traceable evaluation resource.

Each manifest entry contains the following fields: `generator`, `checkpoint`, `prompt_id`, `input_type`, `shotir_spec`, `shotir_to_prompt_converted_text`, `reference_frame_id`, `seed`, `resolution`, `fps`, `num_frames`, `sampling_steps`, `guidance_scale`, `output_path`, `screening_status`, `verified_codes`, `evidence_spans`, and `annotator_ids`. The `screening_status` field takes one of four values: `generated`, `retained_candidate`, `verified`, or `excluded_ambiguous`. Only entries marked as `verified` are included in the reported real-failure evaluation.

```
{
  "video_id": "real_wan_00037",
  "generator": "Wan",
  "checkpoint": "Wan2.1-T2V-14B",
  "prompt_id": "shotir_prompt_037",
  "input_type": "text / ShotIR-to-prompt",
  "shotir_spec": {
    "shot_id": "S037",
    "duration_sec": 8.0,
    "characters": [
      {
        "id": "person_A",
        "description": "young man wearing a blue hoodie",
        "identity_anchor": "ref_person_A_004"
      }
    ]
  },
}
```

**Table 33:** Patch-to-generator adapter details. A repair plan is counted as adapter-executable only when it can be automatically translated into target-generator or editor inputs without manual rewriting. L0/L1 rows denote abstention or prompt-rewrite instructions rather than executable patches.

| Repair-plan action     | ShotIR field              | Counted as | Adapter availability                                      | Concrete generator/editor input                                  | Fallback                                    |
|------------------------|---------------------------|------------|-----------------------------------------------------------|------------------------------------------------------------------|---------------------------------------------|
| inject_reference_frame | character.identity_anchor | L2         | L2 available for SVD; L1 partial for text-to-video models | reference image input + identity prompt strengthening            | prompt-only repair plan                     |
| set_camera_movement    | camera.movement           | L1         | L1 partial                                                | camera-control prompt rewrite; trajectory adapter when supported | no_op_uncertain if no camera control exists |
| set_camera_angle       | camera.viewpoint          | L1         | L1 partial                                                | viewpoint prompt rewrite; optional camera-control token          | repair suggestion only                      |
| inject_prop            | props_required            | L1 / L2    | L1 available; L2 partial with mask/reference input        | prop phrase insertion + optional reference or mask               | prompt edit                                 |
| reposition_prop        | props.region              | L1         | L1 partial                                                | region-aware prompt rewrite or mask-conditioned edit             | prompt-only repair plan                     |
| regenerate_segment     | temporal span             | L2 / L3    | L2 available when segment rerendering is supported        | segment-level rerender or temporal inpainting call               | full-clip rerender                          |
| stabilize_camera       | camera.constraint         | L1 / L2    | L1 partial                                                | no-shake constraint or stabilization post-process                | repair suggestion only                      |
| normalize_frame_rate   | motion.continuity         | L2         | L2 partial                                                | frame interpolation or temporal smoothing post-process           | regenerate_segment                          |
| apply_enhancement      | style.quality             | L2         | L2 available                                              | enhancement / deblocking post-process                            | prompt edit for higher quality              |
| fix_background_anchor  | scene.background_anchor   | L1         | L1 partial                                                | background-anchor prompt strengthening; optional reference image | prompt-only repair plan                     |
| no_op_uncertain        | –                         | L0         | L0 available                                              | abstain from concrete repair                                     | no action                                   |

**Table 34:** Blind human repair evaluation rubric for patch usefulness.


| Score | Criterion                                                                               |
|-------|-----------------------------------------------------------------------------------------|
| 1     | The patch does not address the target failure.                                          |
| 2     | The patch is related to the failure but is not actionable or is largely insufficient.   |
| 3     | The patch is actionable but incomplete or requires substantial manual interpretation.   |
| 4     | The patch directly addresses the target failure with minor ambiguity.                   |
| 5     | The patch directly resolves the target failure and preserves surrounding valid content. |

```

702 "required_props": [
703   {
704     "id": "red_water_bottle",
705     "description": "red water bottle held in the right hand",
706     "expected_presence": "visible throughout the shot"
707   }
708 ],
709 "camera": {
710   "shot_type": "medium shot",
711   "movement": "slow left-to-right tracking",
712   "viewpoint": "front three-quarter view"
713 },
714 "action_order": [
715   "person_A enters the park path",
716   "person_A jogs while holding the red water bottle",
717   "person_A slows down near the bench"
718 ],
719 "style": {
720   "lighting": "natural daylight",
721   "visual_style": "realistic video"
722 },
723 "scene": {
724   "location": "urban park path",
725   "background_anchor": "trees and bench remain stable"
726 }

```





**Figure 14:** Annotation interface used for manual verification in VideoGenDoctor-Bench-v0. The tool exposes a sample navigator, side-by-side review of the original and perturbed videos, temporal evidence marking, structured failure verification cards, and free-form notes for adjudication.


**Table 35:** Error analysis on ambiguous and real-normal samples. The table focuses on no-op behavior and wrong-target repair planning. Best values are boldfaced for interpretable comparison columns; the no\_op\_uncertain rate is reported descriptively rather than bold-ranked.

| Method                        | no_op_uncertain rate | Correct no-op rate | Missed necessary repair | Wrong-target patch | Over-confident patch |
|-------------------------------|----------------------|--------------------|-------------------------|--------------------|----------------------|
| Score-only                    | 0.286                | 0.524              | 0.119                   | 0.171              | 0.321                |
| VLM structured report         | 0.219                | 0.476              | <b>0.096</b>            | 0.229              | 0.396                |
| VLM temporal proposal + patch | 0.257                | 0.518              | 0.108                   | 0.200              | 0.350                |
| VideoGenDoctor-diagnosis      | 0.552                | <b>0.738</b>       | 0.142                   | 0.095              | 0.181                |
| VideoGenDoctor-full           | 0.514                | 0.724              | 0.156                   | <b>0.081</b>       | <b>0.164</b>         |

```

727     },
728     "shotir_to_prompt_converted_text": "A realistic 8-second medium shot of a young man wearing a b
729     "reference_frame_id": null,
730     "seed": 18473,
731     "resolution": "768x432",
732     "fps": 16,
733     "num_frames": 128,
734     "sampling_steps": 50,
735     "guidance_scale": 7.5,
736     "output_path": "real_failures/wan/real_wan_00037.mp4",
737     "screening_status": "verified",
738     "verified_codes": [
739         "AL_PROP_MISSING",
740         "CA_MOVE_WRONG",

```



**Figure 15:** Example VideoGenDoctor diagnostic report for a perturbed video. The report summarizes global scores, ranked failure records with confidence and temporal spans, representative keyframes, patch hints derived from the retained failure codes, and segment-level evidence used for localized review and repair.

**Table 36:** Qualitative case-study summary. Cases are sampled from controlled, real-failure, and ambiguous stress subsets.

| Case type                          | #Cases | Majority-judged successful / appropriate | Typical outcome                               |
|------------------------------------|--------|------------------------------------------|-----------------------------------------------|
| Successful structured repair plan  | 24     | 18                                       | localized repair plan resolves target failure |
| Ambiguous case with abstention     | 24     | 16                                       | no_op_uncertain avoids over-repair            |
| Over-repair / wrong-target failure | 24     | 7                                        | concrete patch targets wrong object or span   |

```

741     "MO_SEGMENT_BREAK"
742 ],
743     "evidence_spans": [
744         {
745             "failure_code": "AL_PROP_MISSING",
746             "target": "red_water_bottle",
747             "start_sec": 3.25,
748             "end_sec": 6.75,
749             "evidence_keyframes": [
750                 "frame_05203.25s",
751                 "frame_08005.00s",
752                 "frame_10806.75s"
753             ],

```

**Table 37:** Leave-one-code-family-out rule ablation. For the held-out family, dedicated family-specific rules are disabled; only generic embedding drift, optical-flow anomaly, scene-change peaks, and structured VLM verification are retained.

| Held-out family | Macro-F1 | tIoU@0.5 | Top-3 hit | Main degradation source                    |
|-----------------|----------|----------|-----------|--------------------------------------------|
| Identity        | 0.782    | 0.579    | 0.536     | face/body-specific anchors removed         |
| Camera          | 0.754    | 0.546    | 0.501     | camera-flow templates disabled             |
| Motion          | 0.721    | 0.512    | 0.468     | jitter / segment-break rules disabled      |
| Alignment       | 0.801    | 0.603    | 0.558     | prop detector and placement rules disabled |
| Style           | 0.836    | 0.641    | 0.590     | compression/style detectors disabled       |
| Scene           | 0.748    | 0.528    | 0.487     | background-anchor rules disabled           |
| Average         | 0.774    | 0.568    | 0.523     | –                                          |

**Table 38:** Generator configuration metadata for the real-failure and real-normal subsets. Hosted or version-changing endpoints are logged with access date and endpoint identifier.

| Generator              | Checkpoint / endpoint | fps | Frames | Resolution | Steps | Guidance | Scheduler      | Seed range  | Reference source / prompt rule     |
|------------------------|-----------------------|-----|--------|------------|-------|----------|----------------|-------------|------------------------------------|
| CogVideoX              | CogVideoX1.5-5B       | 16  | 128    | 768 × 432  | 50    | 6.0      | DPMSolver      | 10000–10059 | ShotIR-to-prompt template v1       |
| Wan                    | Wan2.1-T2V-14B        | 16  | 128    | 768 × 432  | 50    | 5.0      | FlowMatchEuler | 18000–18059 | ShotIR-to-prompt template v1       |
| Stable Video Diffusion | SVD-XT-1.1            | 14  | 112    | 768 × 432  | 40    | 7.5      | EulerDiscrete  | 26000–26059 | reference frame from ShotIR anchor |
| HunyuanVideo           | HunyuanVideo-1.5      | 16  | 128    | 768 × 432  | 50    | 7.0      | FlowMatchEuler | 34000–34059 | multi-shot prompt template v1      |

```

754     "annotator_confidence": 0.91,
755     "notes": "The specified red water bottle is not visible after the character begins jogging
756 },
757 {
758     "failure_code": "CA_MOVE_WRONG",
759     "target": "camera_movement",
760     "start_sec": 0.00,
761     "end_sec": 4.00,
762     "evidence_keyframes": [
763         "frame_000@0.00s",
764         "frame_032@2.00s",
765         "frame_064@4.00s"
766     ],
767     "annotator_confidence": 0.84,
768     "notes": "The camera remains mostly static rather than performing the requested left-to-right
769 },
770 {
771     "failure_code": "MO_SEGMENT_BREAK",
772     "target": "temporal_continuity",
773     "start_sec": 5.50,
774     "end_sec": 5.94,
775     "evidence_keyframes": [
776         "frame_088@5.50s",
777         "frame_095@5.94s"
778     ],
779     "annotator_confidence": 0.78,
780     "notes": "A visible discontinuity occurs near the transition to the final action."
781 }
782 ],
783 "annotator_ids": [
784     "ann_02",
785     "ann_05"
786 ]
787 }

```

For privacy and release safety, file paths may be anonymized, but the manifest must preserve all fields required to reproduce the generation configuration, reconstruct the ShotIR-to-prompt conversion, and audit the verified evidence spans. When a generator does not expose a configuration field such as `guidance_scale`, the field is set to null rather than omitted.

## NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

**The checklist answers are an integral part of your paper submission.** They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”,**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction (Section 1) explicitly state the scope—a controlled diagnostic fixture, not an open-world benchmark—and the four contributions match the experimental results reported in Sections 4.3–4.7. All claims are bounded by the limitations discussed in Section 5.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

840 Question: Does the paper discuss the limitations of the work performed by the authors?

841 Answer: [Yes]

842 Justification: Section 5 includes a dedicated “Limitations” paragraph that discusses the controlled-  
843 fixture scope, failure-enriched sampling, verifier dependency, partial taxonomy coverage, anno-  
844 tation scale, and generator/domain/length coverage gaps. The “When does it fail?” paragraph  
845 further analyzes specific failure modes.

846 Guidelines:

- 847 • The answer [N/A] means that the paper has no limitation while the answer [No] means that  
848 the paper has limitations, but those are not discussed in the paper.
- 849 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 850 • The paper should point out any strong assumptions and how robust the results are to vi-  
851 olations of these assumptions (e.g., independence assumptions, noiseless settings, model  
852 well-specification, asymptotic approximations only holding locally). The authors should  
853 reflect on how these assumptions might be violated in practice and what the implications  
854 would be.
- 855 • The authors should reflect on the scope of the claims made, e.g., if the approach was only  
856 tested on a few datasets or with a few runs. In general, empirical results often depend on  
857 implicit assumptions, which should be articulated.
- 858 • The authors should reflect on the factors that influence the performance of the approach. For  
859 example, a facial recognition algorithm may perform poorly when image resolution is low or  
860 images are taken in low lighting. Or a speech-to-text system might not be used reliably to  
861 provide closed captions for online lectures because it fails to handle technical jargon.
- 862 • The authors should discuss the computational efficiency of the proposed algorithms and how  
863 they scale with dataset size.
- 864 • If applicable, the authors should discuss possible limitations of their approach to address  
865 problems of privacy and fairness.
- 866 • While the authors might fear that complete honesty about limitations might be used by review-  
867 ers as grounds for rejection, a worse outcome might be that reviewers discover limitations that  
868 aren’t acknowledged in the paper. The authors should use their best judgment and recognize  
869 that individual actions in favor of transparency play an important role in developing norms  
870 that preserve the integrity of the community. Reviewers will be specifically instructed to not  
871 penalize honesty concerning limitations.

872 **3. Theory assumptions and proofs**

873 Question: For each theoretical result, does the paper provide the full set of assumptions and a  
874 complete (and correct) proof?

875 Answer: [N/A]

876 Justification: This paper proposes a diagnosis-and-repair framework evaluated empirically. It  
877 does not contain theoretical results, theorems, or proofs.

878 Guidelines:

- 879 • The answer [N/A] means that the paper does not include theoretical results.
- 880 • All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- 881 • All assumptions should be clearly stated or referenced in the statement of any theorems.
- 882 • The proofs can either appear in the main paper or the supplemental material, but if they appear  
883 in the supplemental material, the authors are encouraged to provide a short proof sketch to  
884 provide intuition.
- 885 • Inversely, any informal proof provided in the core of the paper should be complemented by  
886 formal proofs provided in appendix or supplemental material.
- 887 • Theorems and Lemmas that the proof relies upon should be properly referenced.

888 **4. Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The acknowledgments section describes the release package (reference implementation, taxonomy, patch compiler, generation scripts, annotations, evaluation scripts, prompt templates). The main experimental setup is detailed in Section 4.1 and Appendix B, including dataset splits, metrics, and evaluation protocols.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The reference implementation and all generated data (VideoGenDoctor-Bench-v0, annotations, prediction logs) will be released via GitHub under an MIT license upon acceptance. An anonymized repository is included in the supplementary material. Reproduction requires only CPU for table recomputation; GPU is needed only for full feature extraction and VLM judging.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 4.1 and Appendix B specify the controlled-fixture construction (18 clips, 6 domain groups, 9 perturbation operators, 2 seeds), the real-failure sampling protocol (4 generators, 60 failures each, 480 raw  $\rightarrow$  240 retained), feature extractors (OpenCLIP, RAFT/Farneback, InsightFace, YOLOv8), default hyperparameters ( $K = 3$ ,  $\alpha = 0.6$ ,  $\tau = 0.5$ ), and all VLM model versions.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All main result tables report video-level bootstrap 95% confidence intervals. Section 4.7 reports paired bootstrap comparisons with 95% CIs for  $\Delta$  Human Pass@2, explicitly treating estimates with CIs crossing zero as inconclusive. The bootstrap is computed over videos with standard resampling.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.



- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The break-even analysis in Section 5 reports per-video diagnosis cost (\$0.24) and per-video re-rendering cost (\$0.52). Feature extraction uses a single consumer GPU; CPU-only recomputation from serialized predictions completes in seconds. VLM judging uses hosted GPT-4V endpoints.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Section 7 addresses ethical considerations including misuse risk (misleading media), provenance tracking, watermarking, and human review. The annotation protocol was IRB-approved. No human-subject data or PII was collected. All source clips are from permissively licensed public datasets.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 7 discusses positive impact (improved controllability and quality of generated video) and negative impact (potential misuse to strengthen misleading media). Mitigation strategies are suggested: provenance tracking, watermarking, human review, and access control in sensitive deployments.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.



- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [\[Yes\]](#)

Justification: VideoGenDoctor is a diagnostic tool, not a generative model—it does not release any image/video generator that could be directly misused. The released benchmark videos are perturbed diagnostic clips, not photorealistic deepfakes. The paper recommends pairing automated repair with provenance tracking, watermarking, and human review in sensitive deployments (Section 7). The failure taxonomy exposes configurable thresholds for human override.

Guidelines:

- The answer [\[N/A\]](#) means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: All existing assets (CogVideoX, Wan, Stable Video Diffusion, HunyuanVideo, OpenCLIP, RAFT, Farneback, InsightFace, YOLOv8, GPT-4V) are cited with version identifiers. The 18 clean source clips are from publicly available video datasets under permissive licences; sources are listed in the construction manifest (Appendix B).

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

### 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: VideoGenDoctor-Bench-v0 (324 controlled videos with 2,016 evidence spans, plus 240 real-failure videos) and the reference implementation will be released on GitHub under an MIT license. Documentation includes the taxonomy schema (Appendix C), patch compiler specification (Appendix E), annotation guide (Appendix F), VLM baseline prompt templates (Appendix H), and a comprehensive README.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [Yes]

Justification: The annotation protocol involved human annotators from the authors’ institution. Appendix F summarizes the annotation guide, and the annotation interface is shown in Figure 14. Annotators reviewed perturbed videos with specifications and marked failure codes, evidence spans, and keyframes. A dual-annotated reliability subset (120 videos, 768 spans) is reported. Three-annotator stability analysis is in Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [Yes]

1143 Justification: The annotation protocol was reviewed and approved by the authors’ institutional  
1144 review board, as stated in the acknowledgments. The annotation task involved reviewing publicly  
1145 available video clips against textual specifications; participants faced no risks beyond those of  
1146 ordinary screen-based work.

1147 Guidelines:

- 1148 • The answer [N/A] means that the paper does not involve crowdsourcing nor research with  
1149 human subjects.
- 1150 • Depending on the country in which research is conducted, IRB approval (or equivalent) may  
1151 be required for any human subjects research. If you obtained IRB approval, you should  
1152 clearly state this in the paper.
- 1153 • We recognize that the procedures for this may vary significantly between institutions and  
1154 locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines  
1155 for their institution.
- 1156 • For initial submissions, do not include any information that would break anonymity (if  
1157 applicable), such as the institution conducting the review.

## 1158 16. Declaration of LLM usage

1159 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-  
1160 standard component of the core methods in this research? Note that if the LLM is used only for  
1161 writing, editing, or formatting purposes and does *not* impact the core methodology, scientific  
1162 rigor, or originality of the research, declaration is not required.

1163 Answer: [N/A]

1164 Justification: VLM-based judges (GPT-4V, Qwen2-VL) are used as baseline comparison methods  
1165 and as an optional verification stage, not as an important, original, or non-standard component  
1166 of the core VideoGenDoctor method. The Stage-1 rule engine and patch compiler operate  
1167 independently of LLMs. LLMs were not used to draft the paper’s core scientific argument.

1168 Guidelines:

- 1169 • The answer [N/A] means that the core method development in this research does not involve  
1170 LLMs as any important, original, or non-standard components.
- 1171 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not be  
1172 described.